

```
1 # Please do not change this cell because some hidden tests might depend on it.
2 import os
3
4 # Otter grader does not handle ! commands well, so we define and use our
5 # own function to execute shell commands.
6 def shell(commands, warn=True):
7     """Executes the string `commands` as a sequence of shell commands.
8
9     Prints the result to stdout and returns the exit status.
10    Provides a printed warning on non-zero exit status unless `warn`
11    flag is unset.
12    """
13    file = os.popen(commands)
14    print(file.read().rstrip('\n'))
15    exit_status = file.close()
16    if warn and exit_status != None:
17        print(f"Completed with errors. Exit status: {exit_status}\n")
18    return exit_status
19
20 shell("""
21 ls requirements.txt >/dev/null 2>&1
22 if [ ! $? = 0 ]; then
23     rm -rf .tmp
24     git clone https://github.com/cs236299-2020/project4.git .tmp
25     mv .tmp/requirements.txt .
26     rm -rf .tmp
27 fi
28 pip install -q -r requirements.txt
29 """)
```

```
1 # Initialize Otter
2 import otter
3 grader = otter.Notebook()
```

```
%%latex \newcommand{\vect}[1]{\mathbf{#1}} \newcommand{\cnt}[1]{\sharp{#1}} \newcommand{\argmax}[1]{\underset{#1}{\operatorname{argmax}}} \newcommand{\softmax}{\operatorname{softmax}} \newcommand{\Prob}{\Pr} \newcommand{\given}{\mid}
```

▼ Project 4: Semantic parsing for question answering

The goal of semantic parsing is to convert natural language utterances to a meaning representation such as a *logical form* expression or a *SQL query*. In the previous project segment, you built a parsing system to reconstruct parse trees from the natural-language queries in the ATIS dataset. However, that only solves an intermediary task, not the end-user task of obtaining answers to the queries.

In this final project segment, you will go further, building a semantic parsing system to convert the English queries to SQL queries, so that by consulting a database you will be able to answer those questions. You will implement both a rule-based approach and an end-to-end sequence-to-sequence (seq2seq) approach. Both algorithms come with their pros and cons, and by the end of this segment you should have a basic understanding of the characteristics of the rule-based computational linguistic approach and the neural approach.

Goals

1. Build a semantic parsing algorithm to convert text to SQL queries based on the syntactic parse trees from the last project.
2. Build an end-to-end seq2seq system to convert text to SQL.
3. Discuss the pros and cons of the rule-based system and the end-to-end system.
4. (Optional) Use the state-of-the-art pretrained transformers for text-to-SQL conversion.

This will be an extremely challenging project, so we recommend that you start early.

▼ Setup

```
1 import copy
2 import datetime
3 import math
4 import re
5 import sys
6 import warnings
7
8 import nltk
9 import sqlite3
10 import torch
11 import torch.nn as nn
12 import torchtext as tt
13
14 from cryptography.fernet import Fernet
15 from func_timeout import func_set_timeout
16 from torch.nn.utils.rnn import pack_padded_sequence as pack
17 from torch.nn.utils.rnn import pad_packed_sequence as unpack
18 from tqdm import tqdm
19 from transformers import BartTokenizer, BartForConditionalGeneration
```

```
1 # Set random seeds
2 seed = 1234
3 torch.manual_seed(seed)
4
5 # Turn off annoying torchtext warnings about pending deprecations
6 warnings.filterwarnings("ignore", module="torchtext", category=UserWarning)
7
8 # Set timeout for executing SQL
9 TIMEOUT = 3 # seconds
10
11 # GPU check: Set runtime type to use GPU where available
12 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
13 print (device)

cuda
```

```

1 ## Download needed scripts and data
2 source_url = "https://raw.githubusercontent.com/nlp-236299/data/master"
3
4 # Grammar to augment for this segment
5 if not os.path.isfile('data/grammar'):
6     shell(f"wget -nv -N -P data {source_url}/ATIS/grammar_distrib4.crypt")
7
8 # Decrypt the grammar file
9 key = b'bfksTY2BJ5VKKK9xZb1PDDLaGkdu7KCDFYfVePSEfGY='
10 fernet = Fernet(key)
11 with open('./data/grammar_distrib4.crypt', 'rb') as f:
12     restored = Fernet(key).decrypt(f.read())
13 with open('./data/grammar', 'wb') as f:
14     f.write(restored)
15
16 # Download scripts and ATIS database
17 shell(f"""
18     wget -nv -N -P scripts {source_url}/scripts/trees/transform.py
19     wget -nv -N -P data {source_url}/ATIS/atis_sqlite.db
20 """)

```

```

1 # Import downloaded scripts for parsing augmented grammars
2 sys.path.insert(1, './scripts')
3 import transform as xform

```

▼ Semantically augmented grammars

In the first part of this project segment, you'll be implementing a rule-based system for semantic interpretation of sentences. Before jumping into using such a system on the ATIS dataset – we'll get to that soon enough – let's first work with some trivial examples to get things going.

The fundamental idea of rule-based semantic interpretation is the rule of compositionality, that *the meaning of a constituent is a function of the meanings of its immediate subconstituents and the syntactic rule that combined them*. This leads to an infrastructure

for specifying semantic interpretation in which each syntactic rule in a grammar (in our case, a context-free grammar) is associated with a semantic rule that applies to the meanings associated with the elements on the right-hand side of the rule.

Example: arithmetic expressions

As a first example, let's consider an augmented grammar for arithmetic expressions, familiar from lab 3-1. We again use the function `xform.parse_augmented_grammar` to parse the augmented grammar. You can read more about it in the file `scripts/transform.py`.

```

1 arithmetic_grammar, arithmetic_augmentations = xform.parse_augmented_grammar(
2     """
3     ## Sample grammar for arithmetic expressions
4
5     S -> NUM : lambda Num: Num
6         | S OP S : lambda S1, Op, S2: Op(S1, S2)
7
8     OP -> ADD : lambda Op: Op
9         | SUB
10        | MULT
11        | DIV
12
13    NUM -> 'zero' : lambda: 0
14        | 'one' : lambda: 1
15        | 'two' : lambda: 2
16        | 'three' : lambda: 3
17        | 'four' : lambda: 4
18        | 'five' : lambda: 5
19        | 'six' : lambda: 6
20        | 'seven' : lambda: 7
21        | 'eight' : lambda: 8
22        | 'nine' : lambda: 9
23        | 'ten' : lambda: 10
24
25    ADD -> 'plus' | 'added' 'to' : lambda: lambda x, y: x + y
26    SUB -> 'minus' : lambda: lambda x, y: x - y
27    MULT -> 'times' | 'multiplied' 'by' : lambda: lambda x, y: x * y
28    DIV -> 'divided' 'by' : lambda: lambda x, y: x / y
29     """

```

Recall that in this grammar specification format, rules that are not explicitly provided with an augmentation (like all the OP rules after the first OP → ADD) are associated with the textually most recent one (`lambda Op: Op`).

The `parse_augmented_grammar` function returns both an NLTK grammar and a dictionary that maps from productions in the grammar to their associated augmentations. Let's examine the returned grammar.

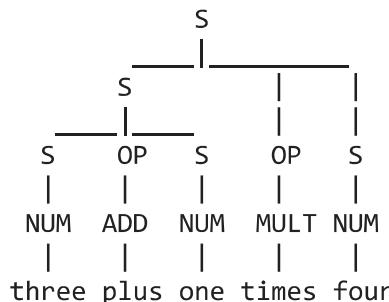
```

1 for production in arithmetic_grammar.productions():
2     print(f"{repr(production):25} {arithmetic_augmentations[production]}")
3
4 S -> NUM                                <function <lambda> at 0x7f9f04cdb730>
5 S -> S OP S                             <function <lambda> at 0x7f9f04cdb620>
6 OP -> ADD                               <function <lambda> at 0x7f9f04cdb378>
7 OP -> SUB                               <function <lambda> at 0x7f9f04cdbd90>
8 OP -> MULT                             <function <lambda> at 0x7f9f04cdb400>
9 OP -> DIV                               <function <lambda> at 0x7f9f04cdb6a8>
10 NUM -> 'zero'                            <function <lambda> at 0x7f9f04cdbf28>
11 NUM -> 'one'                             <function <lambda> at 0x7f9f04cdb2f0>
12 NUM -> 'two'                            <function <lambda> at 0x7f9f04cdbd08>
13 NUM -> 'three'                           <function <lambda> at 0x7f9f04cdbea0>
14 NUM -> 'four'                            <function <lambda> at 0x7f9f04d70bf8>
15 NUM -> 'five'                            <function <lambda> at 0x7f9f04d70488>
16 NUM -> 'six'                             <function <lambda> at 0x7f9f04d70400>
17 NUM -> 'seven'                           <function <lambda> at 0x7f9f04d70378>
18 NUM -> 'eight'                           <function <lambda> at 0x7f9f04d70510>
19 NUM -> 'nine'                            <function <lambda> at 0x7f9f04d702f0>
20 NUM -> 'ten'                             <function <lambda> at 0x7f9f04d70268>
21 ADD -> 'plus'                            <function <lambda> at 0x7f9f04d701e0>
22 ADD -> 'added' 'to'                      <function <lambda> at 0x7f9f04d70158>
23 SUB -> 'minus'                           <function <lambda> at 0x7f9f04d700d0>
24 MULT -> 'times'                           <function <lambda> at 0x7f9f04d70048>
25 MULT -> 'multiplied' 'by'                <function <lambda> at 0x7f9f04d709d8>
26 DIV -> 'divided' 'by'                    <function <lambda> at 0x7f9f04d70c80>
```

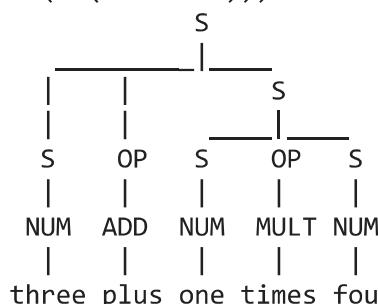
We can parse with the grammar using one of the built-in NLTK parsers.

```
1 arithmetic_parser = nltk.parse.BottomUpChartParser(arithmetic_grammar)
```

```
2 parses = [p for p in arithmetic_parser.parse('three plus one times four'.split())]
3 for parse in parses:
4     parse.pretty_print()
5     print(parse)
```



```
(S
  (S (S (NUM three)) (OP (ADD plus)) (S (NUM one)))
  (OP (MULT times))
  (S (NUM four)))
```



```
(S
  (S (NUM three))
  (OP (ADD plus))
  (S (S (NUM one)) (OP (MULT times)) (S (NUM four))))
```

Now let's turn to the augmentations. They can be arbitrary functions applied to the semantic representations associated with the right-hand-side nonterminals, returning the semantic representation of the left-hand side. To interpret the semantic representation of the entire sentence (at the root of the parse tree), we can use the following pseudo-code:

```
to interpret a tree:
    interpret each of the nonterminal-rooted subtrees
```

```
find the augmentation associated with the root production of the tree
(it should be a function of as many arguments as there are nonterminals on the right-hand side)
return the result of applying the augmentation to the subtree values
```

(The base case of this recursion occurs when the number of nonterminal-rooted subtrees is zero, that is, a rule all of whose right-hand side elements are terminals.)

Suppose we had such a function, call it `interpret`. How would it operate on, for instance, the tree `(S (S (NUM three)) (OP (ADD plus)) (S (NUM one)))`?

```
interpret (S (S (NUM three)) (OP (ADD plus)) (S (NUM one)))
| ->interpret (S (NUM three))
|   | ->interpret (NUM three)
|   |   | ->(no subconstituents to evaluate)
|   |   | ->apply the augmentation for the rule NUM -> three to the empty set of values
|   |   |   | (lambda: 3) () ==> 3
|   |   |   \==> 3
|   |   | ->apply the augmentation for the rule NUM -> NUM to the value 3
|   |   |   | (lambda NUM: NUM)(3) ==> 3
|   |   |   \==> 3
| ->interpret (OP (ADD plus))
|   | ...
|   | \==> lambda x, y: x + y
| ->interpret (S (NUM one))
|   | ...
|   | \==> 1
| ->apply the augmentation for the rule S -> S OP S to the values 3, (lambda x, y: x + y), and 1
|   | (lambda S1, Op, S2: Op(S1, S2))(3, (lambda x, y: x + y), 1) ==> 4
\==> 4
```

Thus, the string "three plus one" is semantically interpreted as the value 4.

Now, all you need to do is to write the `interpret` function. (Lab 4-2 ought to come in handy.)

```

1 #TODO - Write the `interpret` function
2 def interpret(tree, augmentations):
3     if not isinstance(tree, nltk.Tree):
4         return None
5     args = []
6     for child in tree:
7         args.append(interpret(child, augmentations))
8     args = [x for x in args if x is not None]
9     prods = tree.productions()
10    rule = prods[0]
11    augmentation = augmentations[rule]
12    result = augmentation(*args)
13    return result

```

Now we should be able to evaluate the arithmetic example from above.

```
1 interpret(parses[0], arithmetic_augmentations)
```

16

And we can even write a function that parses and interprets a string. We'll have it evaluate each of the possible parses and print the results.

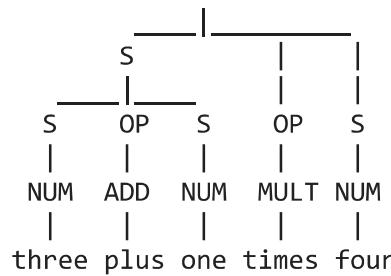
```

1 def parse_and_interpret(string, grammar, augmentations):
2     parser = nltk.parse.BottomUpChartParser(grammar)
3     parses = parser.parse(string.split())
4     for parse in parses:
5         parse.pretty_print()
6         print(parse, "=>", interpret(parse, augmentations))

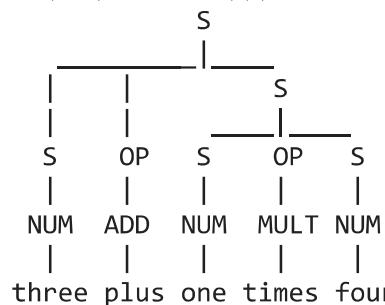
```

```
1 parse_and_interpret("three plus one times four", arithmetic_grammar, arithmetic_augmentations)
```

S



```
(S
  (S (S (NUM three)) (OP (ADD plus)) (S (NUM one)))
  (OP (MULT times))
  (S (NUM four))) ==> 16
```



```
(S
  (S (NUM three))
  (OP (ADD plus))
  (S (S (NUM one)) (OP (MULT times)) (S (NUM four)))) ==> 7
```

Since the string is syntactically ambiguous according to the grammar, it is semantically ambiguous as well.

▼ Some grammar specification conveniences

Before going on, it will be useful to have a few more conveniences in writing augmentations for rules. First, since the augmentations are arbitrary Python expressions, they can be built from and make use of other functions. For instance, you'll notice that many of the augmentations at the leaves of the tree took no arguments and returned a constant. We can define a function `constant` that returns a function that ignores its arguments and returns a particular value.

```
1 def constant(value):
```

```

2     """Return `value`, ignoring any arguments"""
3     return lambda *args: value

```

Similarly, several of the augmentations are functions that just return their first argument. Again, we can define a generic form `first` of such a function:

```

1 def first(*args):
2     """Return the value of the first (and perhaps only) subconstituent,
3         ignoring any others"""
4     return args[0]

```

We can now rewrite the grammar above to take advantage of these shortcuts.

In the call to `parse_augmented_grammar` below, we pass in the `global` environment, extracted via a `globals()` function call, via the named argument `globals`. This allows the `parse_augmented_grammar` function to make use of the global bindings for `constant`, `first`, and the like when evaluating the augmentation expressions to their values. You can check out the code in `transform.py` to see how the passed in `globals` bindings are used. To help understand what's going on, see what happens if you don't include the `globals=globals()`.

```

1 arithmetic_grammar_2, arithmetic_augmentations_2 = xform.parse_augmented_grammar(
2     """
3     ## Sample grammar for arithmetic expressions
4
5     S -> NUM                      : first
6         | S OP S                  : lambda S1, Op, S2: Op(S1, S2)
7
8     OP -> ADD                     : first
9         | SUB
10        | MULT
11        | DIV
12
13    NUM -> 'zero'                 : constant(0)
14        | 'one'                   : constant(1)
15        | 'two'                   : constant(2)

```

```

16     | 'three'           : constant(3)
17     | 'four'            : constant(4)
18     | 'five'             : constant(5)
19     | 'six'              : constant(6)
20     | 'seven'            : constant(7)
21     | 'eight'             : constant(8)
22     | 'nine'              : constant(9)
23     | 'ten'               : constant(10)
24
25 ADD -> 'plus' | 'added' 'to'      : constant(lambda x, y: x + y)
26 SUB -> 'minus'          : constant(lambda x, y: x - y)
27 MULT -> 'times' | 'multiplied' 'by' : constant(lambda x, y: x * y)
28 DIV -> 'divided' 'by'       : constant(lambda x, y: x / y)
29 """
30 globals=globals()

```

Finally, it might make our lives easier to write a template of augmentations whose instantiation depends on the right-hand side of the rule.

We use a reserved keyword `_RHS` to denote the right-hand side of the syntactic rule, which will be replaced by a `list` of the right-hand-side strings. For example, an augmentation `numeric_template(_RHS)` would be as if written as `numeric_template(['zero'])` when the rule is `NUM -> 'zero'`, and `numeric_template(['one'])` when the rule is `NUM -> 'one'`. The details of how this works can be found at [scripts/transform.py](#).

This would allow us to use a single template function, for example,

```

1 def numeric_template(rhs):
2     """Ignore the subphrase meanings and lookup the first right-hand-side symbol
3         as a number"""
4     return constant({'zero':0, 'one':1, 'two':2, 'three':3, 'four':4, 'five':5,
5                     'six':6, 'seven':7, 'eight':8, 'nine':9, 'ten':10}[rhs[0]])

```

and then further simplify the grammar specification:

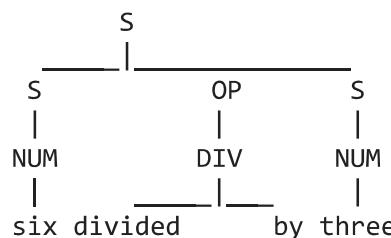
```
1 arithmetic_grammar_3, arithmetic_augmentations_3 = xform.parse_augmented_grammar(
```

```

2     """
3     ## Sample grammar for arithmetic expressions
4
5     S -> NUM : first
6         | S OP S : lambda S1, Op, S2: Op(S1, S2)
7
8     OP -> ADD : first
9         | SUB
10        | MULT
11        | DIV
12
13    NUM -> 'zero' | 'one' | 'two' : numeric_template(_RHS)
14        | 'three' | 'four' | 'five'
15        | 'six' | 'seven' | 'eight'
16        | 'nine' | 'ten'
17
18    ADD -> 'plus' | 'added' 'to' : constant(lambda x, y: x + y)
19    SUB -> 'minus' : constant(lambda x, y: x - y)
20    MULT -> 'times' | 'multiplied' 'by' : constant(lambda x, y: x * y)
21    DIV -> 'divided' 'by' : constant(lambda x, y: x / y)
22    """
23    globals=locals()

```

```
1 parse_and_interpret("six divided by three", arithmetic_grammar_3, arithmetic_augmentations_3)
```



```
(S (S (NUM six)) (OP (DIV divided by)) (S (NUM three))) ==> 2.0
```

▼ Example: Green Eggs and Ham revisited

This stuff is tricky, so it's useful to see more examples before jumping in the deep end. In this simple GEaH fragment grammar, we use a larger set of auxiliary functions to build the augmentations.

```

1 def forward(F, A):
2     """Forward application: Return the application of the first
3         argument to the second"""
4     return F(A)
5
6 def backward(A, F):
7     """Backward application: Return the application of the second
8         argument to the first"""
9     return F(A)
10
11 def second(*args):
12     """Return the value of the second subconstituent, ignoring any others"""
13     return args[1]
14
15 def ignore(*args):
16     """Return `None`, ignoring everything about the constituent. (Good as a
17         placeholder until a better augmentation can be devised.)"""
18     return None

```

Using these, we can build and test the grammar.

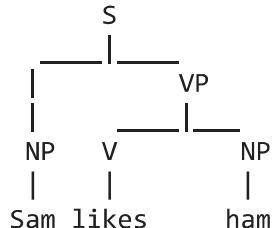
```

1 geah_grammar_spec = """
2     ## Productions
3     S -> NP VP           : backward
4     VP -> V NP            : forward
5
6     ## Lexicon
7     V -> 'likes'          : constant(lambda Object: lambda Subject: f"like({Subject}, {Object})")
8     NP -> 'Sam' | 'sam'   : constant(_RHS[0])
9     NP -> 'ham'
10    NP -> 'eggs'
11 """

```

```
1 geah_grammar, geah_augmentations = xform.parse_augmented_grammar(geah_grammar_spec,
2                                                               globals=globals())
```

```
1 parse_and_interpret("Sam likes ham", geah_grammar, geah_augmentations)
```



```
(S (NP Sam) (VP (V likes) (NP ham))) ==> like(Sam, ham)
```

▼ Semantics of ATIS queries

Now you're in a good position to understand and add augmentations to a more comprehensive grammar, say, one that parses ATIS queries and generates SQL queries.

In preparation for that, we need to load the ATIS data, both NL and SQL queries.

▼ Loading and preprocessing the corpus

To simplify things a bit, we'll only consider ATIS queries whose question type (remember that from project segment 1?) is `flight_id`. We download training, development, and test splits for this subset of the ATIS corpus, including corresponding SQL queries.

```
1 # Acquire the datasets - training, development, and test splits of the
2 # ATIS queries and corresponding SQL queries
3 shell(f"""
4   wget -nv -N -P data {source_url}/ATIS/test_flightid.nl
5   wget -nv -N -P data {source_url}/ATIS/test_flightid.sql
6   wget -nv -N -P data {source_url}/ATIS/dev_flightid.nl
7   wget -nv -N -P data {source_url}/ATIS/dev_flightid.sql
```

```

8 wget -nv -N -P data {source_url}/ATIS/train_flightid.nl
9 wget -nv -N -P data {source_url}/ATIS/train_flightid.sql
10 """")

```

Let's take a look at the data: the NL queries are in .nl files, and the SQL queries are in .sql files.

```

1 shell("head -1 data/dev_flightid.nl")
2 shell("head -1 data/dev_flightid.sql")

what flights are available tomorrow from denver to philadelphia
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_se

```

▼ Corpus preprocessing

We'll use `torchtext` to process the data. We use two `Fields`: `SRC` for the questions, and `TGT` for the SQL queries. We'll use the tokenizer from project segment 3.

```

1 ## Tokenizer
2 tokenizer = nltk.tokenize.RegexpTokenizer('\d+|st\.|[\w-]+|\$[\d\.]+|\s+')
3 def tokenize(string):
4     return tokenizer.tokenize(string.lower())
5
6 ## Demonstrating the tokenizer
7 ## Note especially the handling of `11pm` and hyphenated words.
8 print(tokenize("Are there any first-class flights from St. Louis at 11pm for less than $3.50?"))

```

```
[ 'are', 'there', 'any', 'first-class', 'flights', 'from', 'st.', 'louis', 'at', '11', 'pm', 'for', 'less', 'than', '$3
```

```

1 SRC = tt.data.Field(include_lengths=True,           # include lengths
2                           batch_first=False,          # batches will be max_len x batch_size

```

```

3             tokenize=tokenize,           # use our tokenizer
4         )
5 TGT = tt.data.Field(include_lengths=False,
6                 batch_first=False,      # batches will be max_len x batch_size
7                 tokenize=lambda x: x.split(), # use split to tokenize
8                 init_token="<bos>",       # prepend <bos>
9                 eos_token="<eos>")      # append <eos>
10 fields = [('src', SRC), ('tgt', TGT)]

```

Note that we specified `batch_first=False` (as in lab 4-3), so that the returned batched tensors would be of size `max_length x batch_size`, which facilitates seq2seq implementation.

Now, we load the data using `torchtext`. We use the `TranslationDataset` class here because our task is essentially a translation task: "translating" questions into the corresponding SQL queries. Therefore, we also refer to the questions as the *source side* (`SRC`) and the SQL queries as the *target side* (`TGT`).

```

1 # Make splits for data
2 train_data, val_data, test_data = tt.datasets.TranslationDataset.splits(
3     ('_flightid.nl', '_flightid.sql'), fields, path='./data/',
4     train='train', validation='dev', test='test')
5
6 MIN_FREQ = 3
7 SRC.build_vocab(train_data.src, min_freq=MIN_FREQ)
8 TGT.build_vocab(train_data.tgt, min_freq=MIN_FREQ)
9
10 print(f"Size of English vocab: {len(SRC.vocab)}")
11 print(f"Most common English words: {SRC.vocab.freqs.most_common(10)}\n")
12
13 print(f"Size of SQL vocab: {len(TGT.vocab)}")
14 print(f"Most common SQL words: {TGT.vocab.freqs.most_common(10)}\n")
15
16 print(f"Index for start of sequence token: {TGT.vocab.stoi[TGT.init_token]}")
17 print(f"Index for end of sequence token: {TGT.vocab.stoi[TGT.eos_token]}")

Size of English vocab: 421
Most common English words: [('to', 3478), ('from', 3019), ('flights', 2094), ('the', 1550), ('on', 1230), ('me', 973),

```

```
Size of SQL vocab: 392
Most common SQL words: [('=', 38876), ('AND', 36564), (',', 22772), ('airport_service', 8314), ('city', 8313), ('(', 6
Index for start of sequence token: 2
Index for end of sequence token: 3
```

Next, we batch our data to facilitate processing on a GPU. Batching is a bit tricky because the source and target will typically be of different lengths. Fortunately, `torchtext` allows us to pass in a `sort_key` function. By sorting on length, we can minimize the amount of padding on the source side, but since there is still some padding, we need to handle them with `pack` and `unpack` later on in the seq2seq part (as in lab 4-3).

```
1 BATCH_SIZE = 32 # batch size for training/validation
2 TEST_BATCH_SIZE = 1 # batch size for test, we use 1 to make beam search implementation easier
3
4 train_iter, val_iter = tt.data.BucketIterator.splits((train_data, val_data),
5                                         batch_size=BATCH_SIZE,
6                                         device=device,
7                                         repeat=False,
8                                         sort_key=lambda x: len(x.src),
9                                         sort_within_batch=True)
10 test_iter = tt.data.BucketIterator(test_data,
11                                     batch_size=TEST_BATCH_SIZE,
12                                     device=device,
13                                     repeat=False,
14                                     sort=False,
15                                     train=False)
```

Let's look at a single batch from one of these iterators.

```
1 batch = next(iter(train_iter))
2 train_batch_text, train_batch_text_lengths = batch.src
3 print (f"Size of text batch: {train_batch_text.shape}")
4 print (f"Third sentence in batch: {train_batch_text[:, 2]}")
5 print (f"Length of the third sentence in batch: {train_batch_text_lengths[2]}")
```

```
5 print(f'Length of the third sentence in batch: {train_batch_text[:, 2].shape[0]}')
6 print(f"Converted back to string: {' '.join([SRC.vocab.itos[i] for i in train_batch_text[:, 2]]))}")
7
8 train_batch_sql = batch.tgt
9 print(f"Size of sql batch: {train_batch_sql.shape}")
10 print(f"Third SQL in batch: {train_batch_sql[:, 2]}")
11 print(f"Converted back to string: {' '.join([TGT.vocab.itos[i] for i in train_batch_sql[:, 2]]))}")
```

```
Size of text batch: torch.Size([9, 32])
Third sentence in batch: tensor([ 9,  7, 21,  5, 76, 32,  4, 33, 22], device='cuda:0')
```

Length of the third sentence in batch: 9

Converted back to string: show me all the united airlines flights leaving dallas

```
Size of sql batch: torch.Size([123, 32])
```

Third SQL in batch: tensor([-2, -14, 3])

```

    device='cuda:0')

```

Converted back to string: <bos> SELECT DISTINCT flight 1.flight_id FROM flight flight 1 , airport_service airport_serv

Alternatively, we can directly iterate over the raw examples:

```
1 for example in train_iter.dataset[:1]:  
2     train_text_1 = example.src  
3     train_text_1 = ' '.join(train_text_1) # detokenized question  
4     train_sql_1 = ' '.join(example.tgt)    # detokenized sql  
5     print (f"Question: {train_text_1}\n")  
6     print (f"SQL: {train_sql_1}")
```

Question: list all the flights that arrive at general mitchell international from various cities

SOL: SELECT DISTINCT flight 1.flight_id FROM flight flight 1 . airport airport 1 . airport_service airport_service 1 .

▼ Establishing a SQL database for evaluating ATIS queries

The output of our systems will be SQL queries. How should we determine if the generated queries are correct? We can't merely compare against the gold SQL queries, since there are many ways to implement a SQL query that answers any given NL query.

Instead, we will execute the queries – both the predicted SQL query and the gold SQL query – on an actual database, and verify that the returned responses are the same. For that purpose, we need a SQL database server to use. We'll set one up here, using the [Python sqlite3 module](#).

```
1 @func_set_timeout(TIMEOUT)
2 def execute_sql(sql):
3     conn = sqlite3.connect('data/atis_sqlite.db') # establish the DB based on the downloaded data
4     c = conn.cursor()                         # build a "cursor"
5     c.execute(sql)
6     results = list(c.fetchall())
7     c.close()
8     conn.close()
9     return results
```

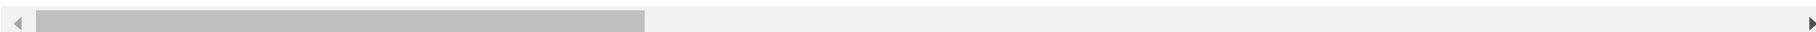
To run a query, we use the cursor's `execute` function, and retrieve the results with `fetchall`. Let's get all the flights that arrive at General Mitchell International – the query `train_sql_1` above. There's a lot, so we'll just print out the first few.

```
1 predicted_ret = execute_sql(train_sql_1)
2
3 print(f"""
4 Executing: {train_sql_1}
5
6 Result: {len(predicted_ret)} entries starting with
7
8 {predicted_ret[:10]}
9 """)
```

Executing: SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport airport_1 , airport_service airport_servi

```
Result: 534 entries starting with
```

```
[(107929,), (107930,), (107931,), (107932,), (107933,), (107934,), (107935,), (107936,), (107937,), (107938,)]
```



For your reference, the SQL database we are using has a database schema described at <https://github.com/jkkummerfeld/text2sql-data/blob/master/data/atis-schema.csv>, and is consistent with the SQL queries provided in the various .sql files loaded above.

▼ Rule-based parsing and interpretation of ATIS queries

First, we will implement a rule-based semantic parser using a grammar like the one you completed in the third project segment. We've placed an initial grammar in the file `data/grammar`. In addition to the helper functions defined above (`constant`, `first`, etc.), it makes use of some other simple functions. We've included those below, but you can (and almost certainly should) augment this set with others that you define as you build out the full set of augmentations.

```
1
2 def upper(term):
3     return ''' + term.upper() + '''
4
5 def weekday(day):
6     return f"flight.flight_days IN (SELECT days.days_code FROM days WHERE days.day_name = '{day.upper()}')"
7
8 def month_name(month):
9     return {'JANUARY' : 1,
10            'FEBRUARY' : 2,
11            'MARCH' : 3,
12            'APRIL' : 4,
13            'MAY' : 5,
14            'JUNE' : 6,
15            'JULY' : 7,
16            'AUGUST' : 8,
17            'SEPTEMBER' : 9,
18            'OCTOBER' : 10,
```

```
19     'NOVEMBER' : 11,
20     'DECEMBER' : 12}[month.upper()]
21
22 def airports_from_airport_name(airport_name):
23     return f"(SELECT airport.airport_code FROM airport WHERE airport.airport_name = {upper(airport_name)})"
24
25 def airports_from_city(city):
26     return f"""
27     (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
28         (SELECT city.city_code FROM city WHERE city.city_name = {upper(city)}))
29 """
30
31
32 def destination_airports(*args):
33     return lambda x : f"""
34     flight.to_airport IN {x}
35 """
36
37 def origin_airports(*args):
38     return lambda x : f"""
39     flight.from_airport IN {x}
40 """
41
42 def between(*args):
43     return f"""
44     flight.from_airport IN {args[0]}
45     AND flight.to_airport IN {args[1]}
46 """
47
48 def airline_pp_brand(*args):
49     return lambda x : f"""
50     flight.airline_code = '{x}'
51 """
52
53 def airline_adj_brand(*args):
54     return f"""
55     flight.airline_code = '{args[0]}'
56 """
57
```

```
58 def null_condition(*args, **kwargs):
59     return 1
60
61 def and_condition(*args):
62     return f"""
63     {args[1]} AND {args[0]}
64 """
65 def s_rule(*args):
66     return f"""
67     SELECT DISTINCT flight.flight_id FROM flight WHERE {args[0]}
68 """
69 def depart_around(time):
70     return f"""
71         flight.departure_time >= {add_delta(miltime(time), -15).strftime('%H%M')}
72         AND flight.departure_time <= {add_delta(miltime(time), 15).strftime('%H%M')}
73     """.strip()
74
75 def arrives_before(*args):
76     return lambda x : f"""
77         flight.arrival_time < {x}
78 """
79
80 def arrives_at(*args):
81     return lambda x : f"""
82         flight.arrival_time == {x}
83 """
84
85 def departs_before(*args):
86     return lambda x : f"""
87         flight.departure_time < {x}
88 """
89
90 def departs_at(*args):
91     return lambda x : f"""
92         flight.departure_time = {x}
93 """
94
95 def departs_after(*args):
96     return lambda x : f"""
```

```

97     flight.departure_time > {x}
98 """
99
100 def arrives_after(*args):
101     return lambda x : f"""
102         flight.arrival_time > {x}
103 """
104
105 def leaves_on(*args):
106     return lambda x : f"""
107     {x}
108 """
109
110
111
112 def add_delta(tme, delta):
113     # transform to a full datetime first
114     return (datetime.datetime.combine(datetime.date.today(), tme) +
115             datetime.timedelta(minutes=delta)).time()
116
117 def militime(minutes):
118     return datetime.time(hour=int(minutes/100), minute=(minutes % 100))

```

We can build a parser with the augmented grammar:

```

1 atis_grammar, atis_augmentations = xform.read_augmented_grammar('data/grammar', globals=globals())
2 atis_parser = nltk.parse.BottomUpChartParser(atis_grammar)

```

We'll define a function to return a parse tree for a string according to the ATIS grammar (if available).

```

1 def parse_tree(sentence):
2     """Parse a sentence and return the parse tree, or None if failure."""
3     try:
4         parses = list(atis_parser.parse(tokenize(sentence)))
5         if len(parses) == 0:
6             return None

```

```

7     else:
8         return parses[0]
9     except:
10        return None

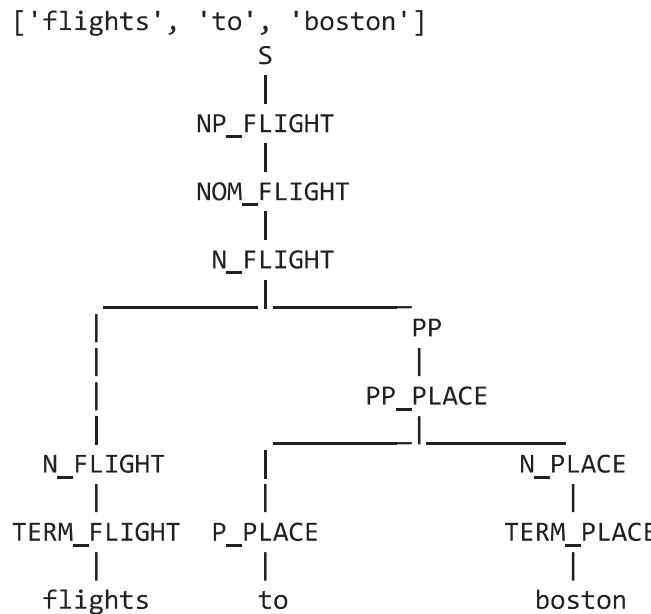
```

and use it to parse the sample query "flights to boston":

```

1 sample_query = "flights to boston"
2 print(tokenize(sample_query))
3 parse_tree(sample_query).pretty_print()

```



We can check the overall coverage of this grammar on the training set by using the `parse_tree` function to determine if a parse is available. The grammar that we provide should get about a 40% coverage of the training set.

```

1 # Check coverage on training set
2 parsed = 0
3 with open("data/train_flightid.nl") as train:
4     for line in train:
5         try:
6             tree = parse_tree(line)
7             parsed += 1
8         except:
9             pass
10    print(parsed)

```

```
4 examples = train.readlines()[:]
5 for sentence in tqdm(examples):
6     if parse_tree(sentence):
7         parsed += 1
8     else:
9         next
10
11 print(f"parsed {parsed} of {len(examples)} ({parsed*100/(len(examples)):.2f}%)")
```

100% |██████████| 3651/3651 [00:14<00:00, 254.52it/s] Parsed 1525 of 3651 (41.77%)

▼ Goal 1: Construct SQL queries from a parse tree and evaluate the results

It's time to turn to the first major part of this project segment, implementing a rule-based semantic parsing system to answer flight-ID-type ATIS queries.

Recall that in rule-based semantic parsing, each syntactic rule is associated with a semantic composition rule. The grammar we've provided has semantic augmentations for some of the low-level phrases – cities, airports, times, airlines – but not the higher level syntactic types. You'll be adding those.

In the ATIS grammar that we provide, as with the earlier toy grammars, the augmentation for a rule with n nonterminals and m terminals on the right-hand side is assumed to be called with n positional arguments (the values for the corresponding children). The `interpret` function you've already defined should therefore work well with this grammar.

Let's run through one way that a semantic derivation might proceed, for the sample query "flights to boston" that we parsed above.

Given a sentence, we first construct its parse tree using the syntactic rules, then compose the corresponding semantic rules bottom-up, until eventually we arrive at the root node with a finished SQL statement. For this query, we will go through what the possible meaning representations for the subconstituents of "flights to boston" might be. But this is just one way of doing things; other ways are possible, and you should feel free to experiment.

Working from bottom up:

1. The `TERM_PLACE` phrase "boston" uses the composition function template `constant(airports_from_city(' '.join(_RHS)))`, which will be instantiated as `constant(airports_from_city(' '.join(['boston'])))` (recall that `_RHS` is replaced by the right-hand side of the rule). The meaning of `TERM_PLACE` will be the SQL snippet

```
SELECT airport_service.airport_code
FROM airport_service
WHERE airport_service.city_code IN
(SELECT city.city_code
FROM city
WHERE city.city_name = "BOSTON")
```

(This query generates a list of all of the airports in Boston.)

2. The `N_PLACE` phrase "boston" can have the same meaning as the `TERM_PLACE`.
3. The `P_PLACE` phrase "to" might be associated with a function that takes a SQL query for a list of airports to a SQL condition (`WHERE body`) that holds of flights that go to one of those airports.
4. The `PP_PLACE` phrase "to boston" might apply the `P_PLACE` meaning to the `TERM_PLACE` meaning, thus generating a SQL condition (`WHERE body`) that holds of flights that go to one of the Boston airports:

```
flight.to_airport IN
(SELECT airport_service.airport_code
FROM airport_service
WHERE airport_service.city_code IN
(SELECT city.city_code
FROM city
WHERE city.city_name = "BOSTON")
```

5. The `PP` phrase "to Boston" can again get its meaning from the `PP_PLACE`.

6. The `TERM_FLIGHT` phrase "flights" might also return a condition on flights, this time the "null condition", represented by the SQL truth value `1`. Ditto for the `N_FLIGHT` phrase "flights".
7. The `N_FLIGHT` phrase "flights to boston" can conjoin the two conditions, yielding the SQL condition

```
flight.to_airport IN
  (SELECT airport_service.airport_code
   FROM airport_service
   WHERE airport_service.city_code IN
     (SELECT city.city_code
      FROM city
      WHERE city.city_name = "BOSTON"))
AND 1
```

which can be inherited by the `NOM_FLIGHT` and `NP_FLIGHT` phrases.

8. The `s` phrase "flights to boston" can use the condition provided by the `NP_FLIGHT` phrase to select all flights satisfying the condition with a SQL query like

```
SELECT DISTINCT flight.flight_id
FROM flight
WHERE flight.to_airport IN
  (SELECT airport_service.airport_code
   FROM airport_service
   WHERE airport_service.city_code IN
     (SELECT city.city_code
      FROM city
      WHERE city.city_name = "BOSTON"))
AND 1
```

This SQL query is then taken to be a representation of the meaning for the NL query "flights to boston", and can be executed against the ATIS database to retrieve the requested flights.

The augmentations that we have provided for the grammar make use of a set of auxiliary functions that we defined above. You should feel free to add your own auxiliary functions that you make use of in the grammar.

▼ Verification on some examples

With a rule-based semantic parsing system, we can generate SQL queries given questions, and then execute those queries on a SQL database to answer the given questions. To evaluate the performance of the system, we compare the returned results against the results of executing the ground truth queries.

We provide a function `verify` to compare the results from our generated SQL to the ground truth SQL. It should be useful for testing individual queries.

```

1 def verify(predicted_sql, gold_sql, silent=True):
2     """
3         Compare the correctness of the generated SQL by executing on the
4         ATIS database and comparing the returned results.
5         Arguments:
6             predicted_sql: the predicted SQL query
7             gold_sql: the reference SQL query to compare against
8             silent: print outputs or not
9         Returns: True if the returned results are the same, otherwise False
10    """
11    # Execute predicted SQL
12    try:
13        predicted_result = execute_sql(predicted_sql)
14    except BaseException as e:
15        if not silent:
16            print(f"predicted sql exec failed: {e}")
17        return False
18    if not silent:
19        print("Predicted DB result:\n\n", predicted_result[:10], "\n")

```

```

20
21 # Execute gold SQL
22 try:
23     gold_result = execute_sql(gold_sql)
24 except BaseException as e:
25     if not silent:
26         print(f"gold sql exec failed: {e}")
27     return False
28 if not silent:
29     print("Gold DB result:\n\n", gold_result[:10], "\n")
30
31 # Verify correctness
32 if gold_result == predicted_result:
33     return True

```

Let's try this methodology on a simple example: "flights from phoenix to milwaukee". we provide it along with the gold SQL query.

```

1 def rule_based_trial(sentence, gold_sql):
2     print("Sentence: ", sentence, "\n")
3     tree = parse_tree(sentence)
4     print("Parse:\n\n", tree, "\n")
5
6     predicted_sql = interpret(tree, atis_augmentations)
7     print("Predicted SQL:\n\n", predicted_sql, "\n")
8
9     if verify(predicted_sql, gold_sql, silent=False):
10        print ('Correct!')
11    else:
12        print ('Incorrect!')

```

```

1 # Example 1
2 example_1 = 'flights from phoenix to milwaukee'
3 gold_sql_1 = """
4     SELECT DISTINCT flight_1.flight_id
5     FROM flight flight_1 ,
6          airport_service airport_service_1 ,
7          city city_1 .

```

```

    ,  

8     airport_service airport_service_2 ,  

9     city city_2  

10 WHERE flight_1.from_airport = airport_service_1.airport_code  

11     AND airport_service_1.city_code = city_1.city_code  

12     AND city_1.city_name = 'PHOENIX'  

13     AND flight_1.to_airport = airport_service_2.airport_code  

14     AND airport_service_2.city_code = city_2.city_code  

15     AND city_2.city_name = 'MILWAUKEE'  

16 """  

17  

18 rule_based_trial(example_1, gold_sql_1)

```

Sentence: flights from phoenix to milwaukee

Parse:

```

(S
 (NP_FLIGHT
  (NOM_FLIGHT
   (N_FLIGHT
    (N_FLIGHT
     (N_FLIGHT (TERM_FLIGHT flights))
     (PP
      (PP_PLACE (P_PLACE from) (N_PLACE (TERM_PLACE phoenix))))
     (PP (PP_PLACE (P_PLACE to) (N_PLACE (TERM_PLACE milwaukee)))))))

```

Predicted SQL:

```

SELECT DISTINCT flight.flight_id FROM flight WHERE

flight.to_airport IN
(SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
(SELECT city.city_code FROM city WHERE city.city_name = "MILWAUKEE"))

```

AND

```

flight.from_airport IN
(SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
(SELECT city.city_code FROM city WHERE city.city_name = "PHOENIX"))

```

AND 1

Predicted DB result:

```
[(108086,), (108087,), (301763,), (301764,), (301765,), (301766,), (302323,), (304881,), (310619,), (310620,)]
```

Gold DB result:

```
[(108086,), (108087,), (301763,), (301764,), (301765,), (301766,), (302323,), (304881,), (310619,), (310620,)]
```

Correct!

To make development faster, we recommend starting with a few examples before running the full evaluation script. We've taken some examples from the ATIS dataset including the gold SQL queries that they provided. Of course, yours (and those of the project segment solution set) may differ.

```
1 # Example 2
2 example_2 = 'i would like a united flight'
3 gold_sql_2 = """
4     SELECT DISTINCT flight_1.flight_id
5     FROM flight flight_1
6     WHERE flight_1.airline_code = 'UA'
7 """
8
9 rule_based_trial(example_2, gold_sql_2)
```

Sentence: i would like a united flight

Parse:

```
(S
  (PREIGNORE
    (PREIGNORESYMBOL i)
  (PREIGNORE
    (PREIGNORESYMBOL would)
  (PREIGNORE
```

```
(PREIGNORESYMBOL like)
(PREIGNORE (PREIGNORESYMBOL a)))))

(NP_FLIGHT
(NOM_FLIGHT
(ADJ (ADJ_AIRLINE (TERM_AIRLINE (TERM_AIRBRAND united))))
(NOM_FLIGHT (N_FLIGHT (TERM_FLIGHT flight)))))
```

Predicted SQL:

```
SELECT DISTINCT flight.flight_id FROM flight WHERE
1 AND
flight.airline_code = 'UA'
```

Predicted DB result:

```
[(100094,), (100099,), (100145,), (100158,), (100164,), (100167,), (100169,), (100203,), (100204,), (100296,)]
```

Gold DB result:

```
[(100094,), (100099,), (100145,), (100158,), (100164,), (100167,), (100169,), (100203,), (100204,), (100296,)]
```

Correct!

```
1 # Example 3
2 example_3 = 'i would like a flight between boston and dallas'
3 gold_sql_3 = """
4     SELECT DISTINCT flight_1.flight_id
5     FROM flight flight_1 ,
6          airport_service airport_service_1 ,
7          city city_1 ,
8          airport_service airport_service_2 ,
9          city city_2
10    WHERE flight_1.from_airport = airport_service_1.airport_code
11        AND airport_service_1.city_code = city_1.city_code
12        AND city_1.city_name = 'BOSTON'
13        AND flight_1.to_airport = airport_service_2.airport_code
14        AND airport_service_2.city_code = city_2.city_code
```

```

15      AND city_2.city_name = 'DALLAS'
16      """
17
18 rule_based_trial(example_3, gold_sql_3)

```

Sentence: i would like a flight between boston and dallas

Parse:

```

(S
(PREIGNORE
  (PREIGNORESYMBOL i)
(PREIGNORE
  (PREIGNORESYMBOL would)
(PREIGNORE
  (PREIGNORESYMBOL like)
  (PREIGNORE (PREIGNORESYMBOL a)))))

(NP_FLIGHT
  (NOM_FLIGHT
    (N_FLIGHT
      (N_FLIGHT (TERM_FLIGHT flight)))
    (PP
      (PP_PLACE
        between
        (N_PLACE (TERM_PLACE boston))
        and
        (N_PLACE (TERM_PLACE dallas)))))))

```

Predicted SQL:

```

SELECT DISTINCT flight.flight_id FROM flight WHERE
flight.from_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
    (SELECT city.city_code FROM city WHERE city.city_name = "BOSTON"))

AND flight.to_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
    (SELECT city.city_code FROM city WHERE city.city_name = "DALLAS"))

AND 1

```

Predicted DB result:

```
[(103171,), (103172,), (103173,), (103174,), (103175,), (103176,), (103177,), (103178,), (103179,), (103180,)]
```

Gold DB result:

```
[(103171,), (103172,), (103173,), (103174,), (103175,), (103176,), (103177,), (103178,), (103179,), (103180,)]
```

Correct!

```

1 # Example 4
2 example_4 = 'show me the united flights from denver to baltimore'
3 gold_sql_4 = """
4     SELECT DISTINCT flight_1.flight_id
5     FROM flight flight_1 ,
6         airport_service airport_service_1 ,
7         city city_1 ,
8         airport_service airport_service_2 ,
9         city city_2
10    WHERE flight_1.airline_code = 'UA'
11        AND ( flight_1.from_airport = airport_service_1.airport_code
12              AND airport_service_1.city_code = city_1.city_code
13              AND city_1.city_name = 'DENVER'
14              AND flight_1.to_airport = airport_service_2.airport_code
15              AND airport_service_2.city_code = city_2.city_code
16              AND city_2.city_name = 'BALTIMORE' )
17
18 """
19
20 rule_based_trial(example_4, gold_sql_4)
```

Sentence: show me the united flights from denver to baltimore

Parse:

```
(S
  (PREIGNORE
    (PREIGNORESYMBOL show))
```

```
(PREIGNORE
  (PREIGNORESYMBOL me)
  (PREIGNORE (PREIGNORESYMBOL the)))
(NP_FLIGHT
  (NOM_FLIGHT
    (ADJ (ADJ_AIRLINE (TERM_AIRLINE (TERM_AIRBRAND united))))
    (NOM_FLIGHT
      (N_FLIGHT
        (N_FLIGHT
          (N_FLIGHT (TERM_FLIGHT flights))
          (PP
            (PP_PLACE (P_PLACE from) (N_PLACE (TERM_PLACE denver))))
          (PP
            (PP_PLACE (P_PLACE to) (N_PLACE (TERM_PLACE baltimore))))))))
```

Predicted SQL:

```
SELECT DISTINCT flight.flight_id FROM flight WHERE

flight.to_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
    (SELECT city.city_code FROM city WHERE city.city_name = "BALTIMORE"))
```

AND

```
flight.from_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
    (SELECT city.city_code FROM city WHERE city.city_name = "DENVER"))
```

AND 1

```
  AND
  flight.airline_code = 'UA'
```

Predicted DB result:

```
[(101231,), (101233,), (305983,)]
```

Gold DB result:

```
[(101231,), (101233,), (305983,)]
```

Correct!

```
1 # Example 5
2 example_5 = 'show flights from cleveland to miami that arrive before 4pm'
3 gold_sql_5 = """
4     SELECT DISTINCT flight_1.flight_id
5     FROM flight flight_1 ,
6         airport_service airport_service_1 ,
7         city city_1 ,
8         airport_service airport_service_2 ,
9         city city_2
10    WHERE flight_1.from_airport = airport_service_1.airport_code
11        AND airport_service_1.city_code = city_1.city_code
12        AND city_1.city_name = 'CLEVELAND'
13        AND ( flight_1.to_airport = airport_service_2.airport_code
14              AND airport_service_2.city_code = city_2.city_code
15              AND city_2.city_name = 'MIAMI'
16              AND flight_1.arrival_time < 1600 )
17 """
18
19 rule_based_trial(example_5, gold_sql_5)
```

Sentence: show flights from cleveland to miami that arrive before 4pm

Parse:

```
(S
  (PREIGNORE (PREIGNORESYMBOL show))
  (NP_FLIGHT
    (NOM_FLIGHT
      (N_FLIGHT
        (N_FLIGHT
          (N_FLIGHT
            (N_FLIGHT (TERM_FLIGHT flights))
            (PP
              (PP_PLACE
```

```

        (P_PLACE from)
        (N_PLACE (TERM_PLACE cleveland))))
    (PP (PP_PLACE (P_PLACE to) (N_PLACE (TERM_PLACE miami))))
    (PP
        (PP_TIME
            (P_TIME that arrive before)
            (NP_TIME (TERM_TIME (TERM_TIME 4) (TERM_TIMEMOD pm)))))))

```

Predicted SQL:

```

SELECT DISTINCT flight.flight_id FROM flight WHERE

    flight.arrival_time < 1600
    AND

    flight.to_airport IN
        (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
            (SELECT city.city_code FROM city WHERE city.city_name = "MIAMI"))

    AND

    flight.from_airport IN
        (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
            (SELECT city.city_code FROM city WHERE city.city_name = "CLEVELAND"))

    AND 1

```

Predicted DB result:

```
[(107698,), (301117,)]
```

Gold DB result:

```
[(107698,), (301117,)]
```

Correct!


```

40
41     AND flight.to_airport IN (SELECT airport_service.airport_code
42                               FROM airport_service
43                               WHERE airport_service.city_code IN (SELECT city.city_code
44                                     FROM city
45                                     WHERE city.city_name = "CHARLOTTE")))
46 """
47
48 rule_based_trial(example_6, gold_sql_6b)

```

Sentence: okay how about a flight on sunday from tampa to charlotte

Parse:

```

(S
(PREIGNORE
  (PREIGNORESYMBOL okay)
(PREIGNORE
  (PREIGNORESYMBOL how)
(PREIGNORE
  (PREIGNORESYMBOL about)
  (PREIGNORE (PREIGNORESYMBOL a)))))

(NP_FLIGHT
  (NOM_FLIGHT
    (N_FLIGHT
      (N_FLIGHT
        (N_FLIGHT
          (TERM_FLIGHT flight))
        (PP
          (PP_DATE (P_DATE on) (NP_DATE (TERM_WEEKDAY sunday))))
          (PP (PP_PLACE (P_PLACE from) (N_PLACE (TERM_PLACE tampa))))
          (PP (PP_PLACE (P_PLACE to) (N_PLACE (TERM_PLACE charlotte)))))))

```

Predicted SQL:

```

SELECT DISTINCT flight.flight_id FROM flight WHERE

flight.to_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
    (SELECT city.city_code FROM city WHERE city.city_name = "CHARLOTTE"))

```

AND

```
flight.from_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
    (SELECT city.city_code FROM city WHERE city.city_name = "TAMPA"))
```

AND

```
flight.flight_days IN (SELECT days.days_code FROM days WHERE days.day_name = 'SUNDAY')
AND 1
```

Predicted DB result:

```
[(101860,), (101861,), (101862,), (101863,), (101864,), (101865,), (305231,)]
```

Gold DB result:

```
[(101860,), (101861,), (101862,), (101863,), (101864,), (101865,), (305231,)]
```

Correct!

```
1 # Example 7
2 example_7 = 'list all flights going from boston to atlanta before 7 am on thursday'
3 gold_sql_7 = """
4   SELECT DISTINCT flight_1.flight_id
5     FROM flight flight_1 ,
6          airport_service airport_service_1 ,
7          city city_1 ,
8          airport_service airport_service_2 ,
9          city city_2 ,
10         days days_1 ,
11        date_day date_day_1
12    WHERE flight_1.from_airport = airport_service_1.airport_code
13      AND airport_service_1.city_code = city_1.city_code
14      AND city_1.city_name = 'BOSTON'
15      AND ( flight_1.to_airport = airport_service_2.airport_code
16            AND airport_service_2.city_code = city_2.city_code
17            AND city_2.city_name = 'ATLANTA'
```

```

18     AND ( flight_1.flight_days = days_1.days_code
19             AND days_1.day_name = date_day_1.day_name
20             AND date_day_1.year = 1991
21             AND date_day_1.month_number = 5
22             AND date_day_1.day_number = 24
23             AND flight_1.departure_time < 700 ) )
24 """
25
26 # Again, the gold answer above used the exact date, as opposed to the
27 # following approach:
28 gold_sql_7b = """
29     SELECT DISTINCT flight.flight_id
30     FROM flight
31     WHERE ((1
32             AND (((1
33                 AND flight.from_airport IN (SELECT airport_service.airport_code
34                     FROM airport_service
35                     WHERE airport_service.city_code IN (SELECT city.city_code
36                         FROM city
37                         WHERE city.city_name = "BOSTON")))
38                 AND flight.to_airport IN (SELECT airport_service.airport_code
39                     FROM airport_service
40                     WHERE airport_service.city_code IN (SELECT city.city_code
41                         FROM city
42                         WHERE city.city_name = "ATLANTA")))
43                 AND flight.departure_time <= 0700)
44                 AND flight.flight_days IN (SELECT days.days_code
45                     FROM days
46                     WHERE days.day_name = 'THURSDAY'))))
47 """
48
49 rule_based_trial(example_7, gold_sql_7b)

```

Parse:

```

(S
(PREIGNORE (PREIGNORESYMBOL list))
(NP_FLIGHT
  /DET z111

```

```
(PCTI a11)
(NOM_FLIGHT
  (N_FLIGHT
    (N_FLIGHT
      (N_FLIGHT
        (N_FLIGHT (TERM_FLIGHT flights))
        (PP
          (PP_PLACE
            (P_PLACE going from)
            (N_PLACE (TERM_PLACE boston))))))
      (PP
        (PP_PLACE (P_PLACE to) (N_PLACE (TERM_PLACE atlanta))))))
    (PP
      (PP_TIME
        (P_TIME before)
        (NP_TIME (TERM_TIME (TERM_TIME 7) (TERM_TIMEMOD am)))))))
  (PP (PP_DATE (P_DATE on) (NP_DATE (TERM_WEEKDAY thursday)))))))
```

Predicted SQL:

```
SELECT DISTINCT flight.flight_id FROM flight WHERE
flight.flight_days IN (SELECT days.days_code FROM days WHERE days.day_name = 'THURSDAY')
AND
  flight.departure_time < 700
AND
  flight.to_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
   (SELECT city.city_code FROM city WHERE city.city_name = "ATLANTA"))
AND
  flight.from_airport IN
  (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
   (SELECT city.city_code FROM city WHERE city.city_name = "BOSTON"))
AND 1
```

Predicted DB result:

```
[(100014,)]
```

Gold DB result:

```
1 # Example 8
2 example_8 = 'list the flights from dallas to san francisco on american airlines'
3 gold_sql_8 = """
4     SELECT DISTINCT flight_1.flight_id
5     FROM flight flight_1 ,
6          airport_service airport_service_1 ,
7          city city_1 ,
8          airport_service airport_service_2 ,
9          city city_2
10    WHERE flight_1.airline_code = 'AA'
11        AND ( flight_1.from_airport = airport_service_1.airport_code
12              AND airport_service_1.city_code = city_1.city_code
13              AND city_1.city_name = 'DALLAS'
14              AND flight_1.to_airport = airport_service_2.airport_code
15              AND airport_service_2.city_code = city_2.city_code
16              AND city_2.city_name = 'SAN FRANCISCO' )
17 """
18
19 rule_based_trial(example_8, gold_sql_8)
```

Sentence: list the flights from dallas to san francisco on american airlines

Parse:

```
(S
  (PREIGNORE
    (PREIGNORESYMBOL list)
    (PREIGNORE (PREIGNORESYMBOL the)))
  (NP_FLIGHT
    (NOM_FLIGHT
```

```
(N_FLIGHT
  (N_FLIGHT
    (N_FLIGHT
      (N_FLIGHT (TERM_FLIGHT flights))
      (PP
        (PP_PLACE (P_PLACE from) (N_PLACE (TERM_PLACE dallas))))
      (PP
        (PP_PLACE
          (P_PLACE to)
          (N_PLACE (TERM_PLACE san francisco))))
      (PP
        (PP_AIRLINE
          (P_AIRLINE on)
          (TERM_AIRLINE
            (TERM_AIRBRAND american)
            (TERM_AIRBRANDTYPE airlines)))))))
```

Predicted SQL:

```
SELECT DISTINCT flight.flight_id FROM flight WHERE
  flight.airline_code = 'AA'
  AND
  flight.to_airport IN
    (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
      (SELECT city.city_code FROM city WHERE city.city_name = "SAN FRANCISCO"))
  AND
  flight.from_airport IN
    (SELECT airport_service.airport_code FROM airport_service WHERE airport_service.city_code IN
      (SELECT city.city_code FROM city WHERE city.city_name = "DALLAS"))
  AND 1
```

Predicted DB result:

```
[ (108452,), (108454,), (108456,), (111083,), (111085,), (111086,), (111090,), (111091,), (111092,), (111094,)]
```

Gold DB result:

```
[ (108452,), (108454,), (108456,), (111083,), (111085,), (111086,), (111090,), (111091,), (111092,), (111094,)]
```

▼ Systematic evaluation on a test set

We can perform a more systematic evaluation by checking the accuracy of the queries on an entire test set for which we have gold queries. The `evaluate` function below does just this, calculating precision, recall, and F1 metrics for the test set. It takes as argument a "predictor" function, which maps token sequences to predicted SQL queries. We've provided a predictor function for the rule-based model in the next cell (and a predictor for the seq2seq system below when we get to that system).

The rule-based system does not generate predictions for all queries; many queries won't parse. The precision and recall metrics take this into account in measuring the efficacy of the method. The recall metric captures what proportion of *all of the test examples* the system generates a correct query. The precision metric captures what proportion of *all of the test examples for which a prediction is generated* the system generates a correct query. (Recall that F1 is just the geometric mean of precision and recall.)

Once you've made some progress on adding augmentations to the grammar, you can evaluate your progress by seeing if the precision and recall have improved. For reference, the solution code achieves precision of about 71% and recall of about 27% for an F1 of 40%.

```
1 def evaluate(predictor, dataset, num_examples=0, silent=True):
2     """Evaluate accuracy of `predictor` by executing predictions on a
3     SQL database and comparing returned results against those of gold queries.
4
5     Arguments:
6         predictor:    a function that maps a token sequence (provided by torchtext)
7                     to a predicted SQL query string
8         dataset:      the dataset of token sequences and gold SQL queries
9         num_examples: number of examples from `dataset` to use; all of
10                    them if 0
11         silent: if set to False, will print out logs
12     Returns: precision, recall, and F1 score
13     """
14
15     # Prepare to count results
```

```
15 if num_examples <= 0:
16     num_examples = len(dataset)
17 example_count = 0
18 predicted_count = 0
19 correct = 0
20 incorrect = 0
21
22 # Process the examples from the dataset
23 for example in tqdm(dataset[:num_examples]):
24     example_count += 1
25     # obtain query SQL
26     predicted_sql = predictor(example.src)
27     if predicted_sql == None:
28         continue
29     predicted_count += 1
30     # obtain gold SQL
31     gold_sql = ' '.join(example.tgt)
32
33     # check that they're compatible
34     if verify(predicted_sql, gold_sql):
35         correct += 1
36     else:
37         incorrect += 1
38
39 # Compute and return precision, recall, F1
40 precision = correct / predicted_count if predicted_count > 0 else 0
41 recall = correct / example_count
42 f1 = (2 * precision * recall) / (precision + recall) if precision + recall > 0 else 0
43 return precision, recall, f1
```

```
1 def rule_based_predictor(tokens):
2     query = ' '.join(tokens)      # detokenized query
3     tree = parse_tree(query)
4     if tree is None:
5         return None
6     try:
7         predicted_sql = interpret(tree, atis_augmentations)
8     except Exception as err:
```

```
9     return None
10    return predicted_sql

1 precision, recall, f1 = evaluate(rule_based_predictor, test_iter.dataset, num_examples=0)
2 print(f"precision: {precision:.2f}")
3 print(f"recall:      {recall:.2f}")
4 print(f"F1:          {f1:.2f}")

100%|██████████| 332/332 [00:13<00:00, 23.74it/s]precision: 0.85
recall: 0.26
F1: 0.40
```

▼ End-to-End Seq2Seq Model

In this part, you will implement a seq2seq model **with attention mechanism** to directly learn the translation from NL query to SQL. You might find labs 4-3 and 4-4 particularly helpful, as the only difference here is that we are using a different dataset.

Note: We recommend using GPUs to train the model in this part (one way to get GPUs is to use [Google Colab](#), by clicking Menu -> Runtime -> Change runtime type -> GPU), as we need to use a very large model to solve the task well. For development we recommend starting with a smaller model and training for only 1 epoch.

▼ Goal 2: Implement a seq2seq model (with attention)

For the sequence-to-sequence model, you need to implement the class `AttnEncoderDecoder`. A reasonable way to proceed is to implement the following methods:

▼ Model

1. `__init__`: an initializer where you create network modules.

2. forward : given source word ids of size (`max_src_len`, `batch_size`) , source lengths of size (`batch_size`) and decoder input target word ids (`max_tgt_len`, `batch_size`), returns logits (`max_tgt_len`, `batch_size`, `V_tgt`). For better modularity you might want to implement it by implementing two functions `forward_encoder` and `forward_decoder` .

Optimization

3. `train_all`: compute loss on training data, compute gradients, and update model parameters to minimize the loss.
4. `evaluate_ppl`: evaluate the current model's perplexity on a given dataset iterator, we use the perplexity value on the validation set to select the best model.

Decoding

5. `predict` : Generates the target sequence given a list of source tokens using beam search decoding. Note that here you can assume the batch size to be 1 for simplicity.

This implementation is essentially building an entire neural seq2seq system, so expect it to be very challenging. The code you write here can also be used for other seq2seq tasks such as machine translation and document summarization

```

1 MAX_T = 15
2
3 class Beam():
4     """
5     Helper class for storing a hypothesis, its score and its decoder hidden state.
6     """
7
8     def __init__(self, decoder_state, tokens, score):
9         self.decoder_state = decoder_state
10        self.tokens = tokens
11        self.score = score
12
13 class BeamSearcher():
14     """
15     Main class for beam search.
16     """
17
18     def __init__(self, model):
19         self.model = model

```

```
20     self.bos_id = model.bos_id
21     self.eos_id = model.eos_id
22     self.padding_id_src = model.padding_id_src
23     self.V = model.V_tgt
24
25     def beam_search(self, src, src_lengths, K, max_T=MAX_T):
26         """
27             Performs beam search decoding.
28             Arguments:
29                 src: src batch of size (max_src_len, 1)
30                 src_lengths: src lengths of size (1)
31                 K: beam size
32                 max_T: max possible target length considered
33             Returns:
34                 a list of token ids and a list of attentions
35         """
36         finished = []
37         all_attns = []
38         # Initialize the beam
39         self.model.eval()
40         # TODO - fill in `memory_bank`, `encoder_final_state`, and `init_beam` below
41
42         (memory_bank, encoder_final_state) = self.model.forward_encoder(src, src_lengths)
43         bos = torch.tensor([[self.model.bos_id]], device=device)
44         init_beam = Beam(encoder_final_state, bos, 0)
45
46         beams = [init_beam]
47
48         with torch.no_grad():
49             for t in range(max_T): # main body of search over time steps
50
51                 # Expand each beam by all possible tokens  $y_{t+1}$ 
52                 all_total_scores = []
53                 for beam in beams:
54                     y_1_to_t, score, decoder_state = beam.tokens, beam.score, beam.decoder_state
55                     y_t = y_1_to_t[-1]
56                     # TODO - finish the code below
57                     # Hint: you might want to use `model.forward_decoder_incrementally`
58                     src_mask = src.ne(self.padding_id_src)
```

```
59
60     logits, decoder_state, attn = self.model.forward_decoder_incrementally(decoder_state, y_t, memory_bank,
61                                         src_mask)
62
63     total_scores = logits
64     all_total_scores.append(total_scores)
65     all_attns.append(attn) # keep attentions for visualization
66     beam.decoder_state = decoder_state # update decoder state in the beam
67     all_total_scores = torch.stack(all_total_scores) # (K, V) when t>0, (1, V) when t=0
68
69     # Find K best next beams
70     # The code below has the same functionality as line 6-12, but is more efficient
71     all_scores_flattened = all_total_scores.view(-1) # K*V when t>0, 1*V when t=0
72     topk_scores, topk_ids = all_scores_flattened.topk(K, 0)
73     beam_ids = topk_ids // self.V
74     next_tokens = topk_ids - beam_ids * self.V
75     new_beams = []
76     for k in range(K):
77         beam_id = beam_ids[k] # which beam it comes from
78         y_t_plus_1 = next_tokens[k] # which y_{t+1}
79         score = topk_scores[k]
80         beam = beams[beam_id]
81         decoder_state = beam.decoder_state
82         y_1_to_t = beam.tokens
83         # TODO
84         y_t = torch.cat((y_1_to_t, y_t_plus_1.view(1, 1)))
85         new_beam = Beam(decoder_state, y_t, beam.score + score)
86         new_beams.append(new_beam)
87     beams = new_beams
88
89     # Set aside completed beams
90     # TODO - move completed beams to `finished` (and remove them from `beams`)
91     new_beams = []
92     for beam in beams:
93         y_1_to_t, score, decoder_state = beam.tokens, beam.score, beam.decoder_state
94         y_t = y_1_to_t[-1]
95         if y_t.view(1) == self.model.eos_id:
96             finished.append(beam)
97         else:
```

```
98     new_beams.append(beam)
99     beams = new_beams
100
101    # Break the loop if everything is completed
102    if len(beams) == 0:
103        break
104
105    # Return the best hypothesis
106    if len(finished) > 0:
107        finished = sorted(finished, key=lambda beam: -beam.score)
108        return finished[0].tokens, all_attns
109    else: # when nothing is finished, return an unfinished hypothesis
110        return beams[0].tokens, all_attns
111
112 def attention(batched_Q, batched_K, batched_V, mask=None):
113     """
114     Performs the attention operation and returns the attention matrix
115     `batched_A` and the context matrix `batched_C` using queries
116     `batched_Q`, keys `batched_K`, and values `batched_V`.
117     Arguments:
118         batched_Q: (q_len, bsz, D)
119         batched_K: (k_len, bsz, D)
120         batched_V: (k_len, bsz, D)
121         mask: (bsz, q_len, k_len). An optional boolean mask *disallowing*
122             attention where the mask value is ``False``.
123     Returns:
124         batched_A: the normalized attention scores (bsz, q_len, k_len)
125         batched_C: a tensor of size (q_len, bsz, D).
126     """
127     # Check sizes
128     D = batched_Q.size(-1)
129     bsz = batched_Q.size(1)
130     q_len = batched_Q.size(0)
131     k_len = batched_K.size(0)
132     assert batched_K.size(-1) == D and batched_V.size(-1) == D
133     assert batched_K.size(1) == bsz and batched_V.size(1) == bsz
134     assert batched_V.size(0) == k_len
135     if mask is not None:
136         assert mask.size() == torch.Size([bsz, q_len, k_len])
```

```
137 Q = torch.transpose(batched_Q, dim0=0, dim1=1)
138 K = torch.transpose(batched_K, dim0=0, dim1=1)
139 K = torch.transpose(K, dim0=1, dim1=2)
140 QK = Q @ K
141 if(mask is not None):
142     mask = mask == False
143     QK = QK.masked_fill_(mask, -float("Inf"))
144 batched_A = nn.Softmax(dim=-1)(QK)
145 V = torch.transpose(batched_V, dim0=0, dim1=1)
146 batched_C = batched_A @ V
147 batched_C = torch.transpose(batched_C, dim0=0, dim1=1)
148 # Verify that things sum up to one properly.
149 assert torch.all(torch.isclose(batched_A.sum(-1),
150                         torch.ones(bsz, q_len).to(device)))
151 return batched_A, batched_C
152
153 class AttnEncoderDecoder(nn.Module):
154     def __init__(self, src_field, tgt_field, hidden_size=64, layers=3):
155         """
156             Initializer. Creates network modules and loss function.
157             Arguments:
158                 src_field: src field
159                 tgt_field: tgt field
160                 hidden_size: hidden layer size of both encoder and decoder
161                 layers: number of layers of both encoder and decoder
162         """
163         super(AttnEncoderDecoder, self).__init__()
164         self.src_field = src_field
165         self.tgt_field = tgt_field
166
167         # Keep the vocabulary sizes available
168         self.V_src = len(src_field.vocab.itos)
169         self.V_tgt = len(tgt_field.vocab.itos)
170
171         # Get special word ids
172         self.padding_id_src = src_field.vocab.stoi[src_field.pad_token]
173         self.padding_id_tgt = tgt_field.vocab.stoi[tgt_field.pad_token]
174         self.bos_id = tgt_field.vocab.stoi[tgt_field.init_token]
175         self.eos_id = tgt_field.vocab.stoi[tgt_field.eos_token]
```

```
176
177 # Keep hyper-parameters available
178 self.embedding_size = hidden_size
179 self.hidden_size = hidden_size
180 self.layers = layers
181
182 # Create essential modules
183 self.word_embeddings_src = nn.Embedding(self.V_src, self.embedding_size)
184 self.word_embeddings_tgt = nn.Embedding(self.V_tgt, self.embedding_size)
185
186 # RNN cells
187 self.encoder_rnn = nn.LSTM(
188     input_size      = self.embedding_size,
189     hidden_size    = hidden_size // 2, # to match decoder hidden size
190     num_layers     = layers,
191     bidirectional = True           # bidirectional encoder
192 )
193 self.decoder_rnn = nn.LSTM(
194     input_size      = self.embedding_size,
195     hidden_size    = hidden_size,
196     num_layers     = layers,
197     bidirectional = False          # unidirectional decoder
198 )
199
200 # Final projection layer
201 self.hidden2output = nn.Linear(2*hidden_size, self.V_tgt) # project the concatenation to logits
202
203 # Create loss function
204 self.loss_function = nn.CrossEntropyLoss(reduction='sum',
205                                         ignore_index=self.padding_id_tgt)
206
207 def forward_encoder(self, src, src_lengths):
208 """
209 Encodes source words `src`.
210 Arguments:
211     src: src batch of size (max_src_len, bsz)
212     src_lengths: src lengths of size (bsz)
213 Returns:
214     memory bank: a tensor of size (src len, bsz, hidden size)
```

```
215     (final_state, context): `final_state` is a tuple (h, c) where h/c is of size
216                         (layers, bsz, hidden_size), and `context` is `None`.
217     """
218     # Compute word embeddings
219     src_embeddings = self.word_embeddings_src(src) # max_src_len, bsz, embedding_size
220     src_lengths = src_lengths.tolist()
221     # Deal with paddings
222     packed_src = pack(src_embeddings, src_lengths)
223     #TODO
224     packed_outputs, (h, c) = self.encoder_rnn(packed_src)
225
226     shape = h.shape
227     h = h.view(2, int(shape[0] / 2), shape[1], shape[2])
228     h = torch.cat([h[0], h[1]], dim=2)
229     c = c.view(2, int(shape[0] / 2), shape[1], shape[2])
230     c = torch.cat([c[0], c[1]], dim=2)
231
232     memory_bank, _ = torch.nn.utils.rnn.pad_packed_sequence(packed_outputs)
233     final_state = (h,c)
234     context = None
235     return memory_bank, (final_state, context)
236
237 def forward_decoder(self, encoder_final_state, tgt_in, memory_bank, src_mask):
238     """
239     Decodes based on encoder final state, memory bank, src_mask, and ground truth
240     target words.
241     Arguments:
242         encoder_final_state: (final_state, None) where final_state is the encoder
243                             final state used to initialize decoder. None is the
244                             initial context (there's no previous context at the
245                             first step).
246         tgt_in: a tensor of size (tgt_len, bsz)
247         memory_bank: a tensor of size (src_len, bsz, hidden_size), encoder outputs
248                      at every position
249         src_mask: a tensor of size (src_len, bsz): a boolean tensor, `False` where
250                      src is padding (we disallow decoder to attend to those places).
251     Returns:
252         Logits of size (tgt_len, bsz, V_tgt) (before the softmax operation)
253     """

```

```
254 max_tgt_length = tgt_in.size(0)
255
256 # Initialize decoder state, note that it's a tuple (state, context) here
257 decoder_states = encoder_final_state
258
259 all_logits = []
260 for i in range(max_tgt_length):
261     logits, decoder_states, attn = \
262         self.forward_decoder_incrementally(decoder_states,
263                                             tgt_in[i],
264                                             memory_bank,
265                                             src_mask,
266                                             normalize=False)
267     all_logits.append(logits)          # list of bsz, vocab_tgt
268 all_logits = torch.stack(all_logits, 0) # tgt_len, bsz, vocab_tgt
269 return all_logits
270
271 def forward(self, src, src_lengths, tgt_in):
272     """
273     Performs forward computation, returns logits.
274     Arguments:
275         src: src batch of size (max_src_len, bsz)
276         src_lengths: src lengths of size (bsz)
277         tgt_in: a tensor of size (tgt_len, bsz)
278     """
279     # Create padding mask
280     src_mask = src.ne(self.padding_id_src) # max_src_len, bsz
281     # TODO
282     memory_bank, encoder_final_state = self.forward_encoder(src, src_lengths)
283     logits = self.forward_decoder(encoder_final_state, tgt_in, memory_bank, src_mask)
284     return logits
285
286 def forward_decoder_incrementally(self, prev_decoder_states, tgt_in_onestep,
287                                     memory_bank, src_mask,
288                                     normalize=True):
289     """
290     Forward the decoder for a single step with token `tgt_in_onestep`.
291     This function will be used both in `forward_decoder` and in beam search.
292     Note that bsz can be greater than 1.

```

```
-->      Note that bsz can be greater than 1.
```

293 Arguments:

```
294     prev_decoder_states: a tuple (prev_decoder_state, prev_context). `prev_context`  
295             is `None` for the first step  
296     tgt_in_onestep: a tensor of size (bsz), tokens at one step  
297     memory_bank: a tensor of size (src_len, bsz, hidden_size), encoder outputs  
298             at every position  
299     src_mask: a tensor of size (src_len, bsz): a boolean tensor, `False` where  
300             src is padding (we disallow decoder to attend to those places).  
301     normalize: use log_softmax to normalize or not. Beam search needs to normalize,  
302             while `forward_decoder` does not
```

303 Returns:

```
304     logits: log probabilities for `tgt_in_token` of size (bsz, V_tgt)  
305     decoder_states: (`decoder_state`, `context`) which will be used for the  
306             next incremental update  
307     attn: normalized attention scores at this step (bsz, src_len)  
308 """  
309     prev_decoder_state, prev_context = prev_decoder_states  
310 #TODO  
311     tgt_embeddings = self.word_embeddings_tgt(tgt_in_onestep.view(1, -1)) # max_len, batch_size, embedding_size  
312     if prev_context != None:  
313         tgt_embeddings += prev_context  
314     prev_decoder_state = (prev_decoder_state, prev_context.repeat(self.layers, 1, 1))  
315     decoder_state, (h, c) = self.decoder_rnn(tgt_embeddings, prev_decoder_state)
```

```
316  
317     mask = src_mask.T.unsqueeze(1)  
318     attn, context = attention(decoder_state, memory_bank, memory_bank, mask)
```

```
319  
320     decoder_states = (decoder_state, context)  
321     dec_out_e_ctxt = torch.cat(decoder_states, dim=2)  
322     logits = self.hidden2output(dec_out_e_ctxt).squeeze(0)  
323     attn = attn.squeeze(1)  
324     decoder_states = (decoder_states[0].repeat(self.layers, 1, 1), decoder_states[1])  
325     if normalize:  
326         logits = torch.log_softmax(logits, dim=-1)  
327     return logits, decoder_states, attn
```

```
328  
329     def evaluate_ppl(self, iterator):  
330         """Returns the model's perplexity on a given dataset `iterator`."""
```

```
331         # Switch to eval mode
```

```
331     # Switch to eval mode
332     self.eval()
333     total_loss = 0
334     total_words = 0
335     for batch in iterator:
336         # Input and target
337         src, src_lengths = batch.src
338         tgt = batch.tgt # max_length_sql, bsz
339         tgt_in = tgt[:-1] # remove <eos> for decode input (y_0=<bos>, y_1, y_2)
340         tgt_out = tgt[1:] # remove <bos> as target (y_1, y_2, y_3=<eos>)
341         # Forward to get logits
342         logits = self.forward(src, src_lengths, tgt_in)
343         # Compute cross entropy loss
344         loss = self.loss_function(logits.view(-1, self.V_tgt), tgt_out.view(-1))
345         total_loss += loss.item()
346         total_words += tgt_out.ne(self.padding_id_tgt).float().sum().item()
347     return math.exp(total_loss/total_words)
348
349 def train_all(self, train_iter, val_iter, epochs=10, learning_rate=0.001):
350     """Train the model."""
351     # Switch the module to training mode
352     self.train()
353     # Use Adam to optimize the parameters
354     optim = torch.optim.Adam(self.parameters(), lr=learning_rate)
355     best_validation_ppl = float('inf')
356     best_model = None
357     # Run the optimization for multiple epochs
358     for epoch in range(epochs):
359         total_words = 0
360         total_loss = 0.0
361         for batch in tqdm(train_iter):
362             # Zero the parameter gradients
363             self.zero_grad()
364             # Input and target
365             src, src_lengths = batch.src # text: max_src_length, bsz
366             tgt = batch.tgt # max_tgt_length, bsz
367             tgt_in = tgt[:-1] # Remove <eos> for decode input (y_0=<bos>, y_1, y_2)
368             tgt_out = tgt[1:] # Remove <bos> as target (y_1, y_2, y_3=<eos>)
369             bsz = tgt.size(1)
370             # Run forward pass and compute loss along the way
```

```

370     # Run forward pass and compute loss along the way.
371     logits = self.forward(src, src_lengths, tgt_in)
372     loss = self.loss_function(logits.view(-1, self.V_tgt), tgt_out.view(-1))
373     # Training stats
374     num_tgt_words = tgt_out.ne(self.padding_id_tgt).float().sum().item()
375     total_words += num_tgt_words
376     total_loss += loss.item()
377     # Perform backpropagation
378     loss.div(bsz).backward()
379     optim.step()
380
381     # Evaluate and track improvements on the validation dataset
382     validation_ppl = self.evaluate_ppl(val_iter)
383     self.train()
384     if validation_ppl < best_validation_ppl:
385         best_validation_ppl = validation_ppl
386         self.best_model = copy.deepcopy(self.state_dict())
387     epoch_loss = total_loss / total_words
388     print (f'Epoch: {epoch} Training Perplexity: {math.exp(epoch_loss):.4f} '
389           f'Validation Perplexity: {validation_ppl:.4f}')
390
391     def predict(self, tokens, K, max_T=400):
392         src = [self.src_field.vocab.stoi[token] for token in tokens]
393         src, src_lengths = torch.tensor([src], device=device).T, torch.tensor([len(src)], device=device)
394         # Predict
395         beam_searcher = BeamSearcher(self)
396         prediction, _ = beam_searcher.beam_search(src, src_lengths, K, max_T)
397         # Convert to string
398         prediction = ' '.join([TGT.vocab.itos[token] for token in prediction])
399         prediction = prediction.lstrip('<bos>').rstrip('<eos>').strip()
400     return prediction

```

1

We provide the recommended hyperparameters for the final model in the script below, but you are free to tune the hyperparameters or change any part of the provided code.

For quick debugging, we recommend starting with smaller models (by using a very small `hidden_size`), and only a single epoch. If the model runs smoothly, then you can train the full model on GPUs.

```
1 EPOCHS = 50 # epochs, we recommend starting with a smaller number like 1
2 LEARNING_RATE = 1e-4 # learning rate
3
4 # Instantiate and train classifier
5 model = AttnEncoderDecoder(SRC, TGT,
6     hidden_size      = 1024,
7     layers          = 1,
8 ).to(device)
9
10 model.train_all(train_iter, val_iter, epochs=EPOCHS, learning_rate=LEARNING_RATE)
11 model.load_state_dict(model.best_model)
12
13 # Evaluate model performance, the expected value should be < 1.2
14 print (f'Validation perplexity: {model.evaluate_ppl(val_iter):.3f}')
```

```
100% [██████████] 115/115 [02:04<00:00, 1.09s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 0 Training Perplexity: 11.9775 Validation Perplexity: 2.8618
100% [██████████] 115/115 [02:05<00:00, 1.09s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 1 Training Perplexity: 2.1186 Validation Perplexity: 1.8294
100% [██████████] 115/115 [02:04<00:00, 1.08s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 2 Training Perplexity: 1.6269 Validation Perplexity: 1.5783
100% [██████████] 115/115 [02:07<00:00, 1.11s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 3 Training Perplexity: 1.4676 Validation Perplexity: 1.4687
100% [██████████] 115/115 [02:06<00:00, 1.10s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 4 Training Perplexity: 1.3774 Validation Perplexity: 1.4009
100% [██████████] 115/115 [02:04<00:00, 1.09s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 5 Training Perplexity: 1.3240 Validation Perplexity: 1.3464
100% [██████████] 115/115 [02:02<00:00, 1.07s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 6 Training Perplexity: 1.2810 Validation Perplexity: 1.3293
100% [██████████] 115/115 [02:04<00:00, 1.08s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 7 Training Perplexity: 1.2527 Validation Perplexity: 1.2981
100% [██████████] 115/115 [02:08<00:00, 1.11s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 8 Training Perplexity: 1.2220 Validation Perplexity: 1.2642
100% [██████████] 115/115 [02:05<00:00, 1.09s/it]
 0% [██████████] 0/115 [00:00<?, ?it/s]Epoch: 9 Training Perplexity: 1.2013 Validation Perplexity: 1.2496
```

100%	[██████████]	115/115 [02:06<00:00, 1.10s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 10 Training Perplexity: 1.1780 Validation Perplexity: 1.2286
100%	[██████████]	115/115 [02:03<00:00, 1.07s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 11 Training Perplexity: 1.1599 Validation Perplexity: 1.2079
100%	[██████████]	115/115 [02:02<00:00, 1.07s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 12 Training Perplexity: 1.1428 Validation Perplexity: 1.1938
100%	[██████████]	115/115 [02:05<00:00, 1.09s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 13 Training Perplexity: 1.1330 Validation Perplexity: 1.1876
100%	[██████████]	115/115 [02:05<00:00, 1.09s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 14 Training Perplexity: 1.1236 Validation Perplexity: 1.1808
100%	[██████████]	115/115 [02:05<00:00, 1.10s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 15 Training Perplexity: 1.1147 Validation Perplexity: 1.1779
100%	[██████████]	115/115 [02:04<00:00, 1.08s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 16 Training Perplexity: 1.1024 Validation Perplexity: 1.1698
100%	[██████████]	115/115 [02:08<00:00, 1.11s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 17 Training Perplexity: 1.0963 Validation Perplexity: 1.2118
100%	[██████████]	115/115 [02:05<00:00, 1.09s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 18 Training Perplexity: 1.1141 Validation Perplexity: 1.1815
100%	[██████████]	115/115 [02:06<00:00, 1.10s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 19 Training Perplexity: 1.0908 Validation Perplexity: 1.1606
100%	[██████████]	115/115 [02:06<00:00, 1.10s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 20 Training Perplexity: 1.0836 Validation Perplexity: 1.1671
100%	[██████████]	115/115 [02:04<00:00, 1.08s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 21 Training Perplexity: 1.0760 Validation Perplexity: 1.1560
100%	[██████████]	115/115 [02:05<00:00, 1.09s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 22 Training Perplexity: 1.0683 Validation Perplexity: 1.1533
100%	[██████████]	115/115 [02:02<00:00, 1.06s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 23 Training Perplexity: 1.0634 Validation Perplexity: 1.1578
100%	[██████████]	115/115 [02:04<00:00, 1.08s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 24 Training Perplexity: 1.0592 Validation Perplexity: 1.1545
100%	[██████████]	115/115 [02:05<00:00, 1.09s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 25 Training Perplexity: 1.0555 Validation Perplexity: 1.1494
100%	[██████████]	115/115 [02:03<00:00, 1.07s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 26 Training Perplexity: 1.0506 Validation Perplexity: 1.1510
100%	[██████████]	115/115 [02:02<00:00, 1.07s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 27 Training Perplexity: 1.0531 Validation Perplexity: 1.1521
100%	[██████████]	115/115 [02:06<00:00, 1.10s/it]
0%	[██████████]	0/115 [00:00<?, ?it/s]Epoch: 28 Training Perplexity: 1.0486 Validation Perplexity: 1.1548
100%	[██████████]	115/115 [02:05<00:00, 1.09s/it]

With a trained model, we can convert questions to SQL statements. We recommend making sure that the model can generate at least reasonable results on the examples from before, before evaluating on the full test set.

```
1 torch.save(model.best_model, './best_model')
2 def seq2seq_trial(sentence, gold_sql):
3     print("Sentence: ", sentence, "\n")
4     tokens = tokenize(sentence)
5
6     predicted_sql = model.predict(tokens, K=1, max_T=400)
7     print("Predicted SQL:\n\n", predicted_sql, "\n")
8
9     if verify(predicted_sql, gold_sql, silent=False):
10        print ('Correct!')
11    else:
12        print ('Incorrect!')
```

```
1 seq2seq_trial(example_1, gold_sql_1)
```

Sentence: flights from phoenix to milwaukee

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_s
```

Predicted DB result:

```
[(108086,), (108087,), (301763,), (301764,), (301765,), (301766,), (302323,), (304881,), (310619,), (310620,)]
```

Gold DB result:

```
[(108086,), (108087,), (301763,), (301764,), (301765,), (301766,), (302323,), (304881,), (310619,), (310620,)]
```

Correct!

```
1 seq2seq_trial(example_2, gold_sql_2)
```

Sentence: i would like a united flight

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 WHERE flight
```

Predicted DB result:

```
[(100094,), (100099,), (100699,), (100703,), (100704,), (100705,), (100706,), (101082,), (101083,), (101084,)]
```

Gold DB result:

```
[(100094,), (100099,), (100145,), (100158,), (100164,), (100167,), (100169,), (100203,), (100204,), (100296,)]
```

Incorrect!



```
1 seq2seq_trial(example_3, gold_sql_3)
```

Sentence: i would like a flight between boston and dallas

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_s
```

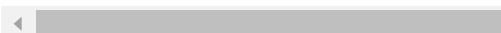
Predicted DB result:

```
[(103171,), (103172,), (103173,), (103174,), (103175,), (103176,), (103177,), (103178,), (103179,), (103180,)]
```

Gold DB result:

```
[(103171,), (103172,), (103173,), (103174,), (103175,), (103176,), (103177,), (103178,), (103179,), (103180,)]
```

Correct!



```
1 seq2seq_trial(example_4, gold_sql_4)
```

Sentence: show me the united flights from denver to baltimore

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_s
```

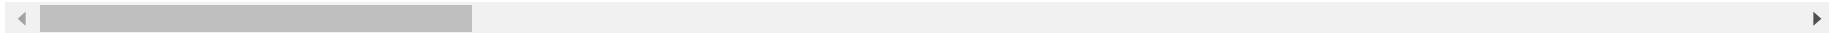
Predicted DB result:

```
[(101231,), (101233,), (305983,)]
```

Gold DB result:

```
[(101231,), (101233,), (305983,)]
```

Correct!



```
1 seq2seq_trial(example_5, gold_sql_5)
```

Sentence: show flights from cleveland to miami that arrive before 4pm

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_s
```

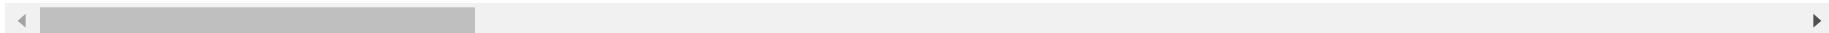
Predicted DB result:

```
[(107698,), (301117,)]
```

Gold DB result:

```
[(107698,), (301117,)]
```

Correct!



```
1 seq2seq_trial(example_6, gold_sql_6b)
```

Sentence: okay how about a flight on sunday from tampa to charlotte

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_s
```

Predicted DB result:

```
[(101860,), (101861,), (101862,), (101863,), (101864,), (101865,), (305231,)]
```

Gold DB result:

```
[(101860,), (101861,), (101862,), (101863,), (101864,), (101865,), (305231,)]
```

Correct!

```
1 seq2seq_trial(example_7, gold_sql_7b)
```

Sentence: list all flights going from boston to atlanta before 7 am on thursday

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_s
```

Predicted DB result:

```
[(100014,)]
```

Gold DB result:

```
[(100014,)]
```

Correct!

```
1 seq2seq_trial(example_8, gold_sql_8)
```

Sentence: list the flights from dallas to san francisco on american airlines

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 , airport_service airport_service_1 , city city_1 , airport_s
```

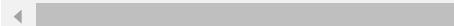
Predicted DB result:

```
[(108452,), (108454,), (108456,), (111083,), (111085,), (111086,), (111090,), (111091,), (111092,), (111094,)]
```

Gold DB result:

```
[(108452,), (108454,), (108456,), (111083,), (111085,), (111086,), (111090,), (111091,), (111092,), (111094,)]
```

Correct!



▼ Evaluation

Now we are ready to run the full evaluation. A proper implementation should reach more than 35% precision/recall/F1.

```
1 def seq2seq_predictor(tokens):
2     prediction = model.predict(tokens, K=1, max_T=400)
3     return prediction
4
5
6 precision, recall, f1 = evaluate(seq2seq_predictor, test_iter.dataset, num_examples=0)
7 print(f"precision: {precision:.2f}")
8 print(f"recall: {recall:.2f}")
9 print(f"F1: {f1:.2f}")
10
11 100%|██████████| 332/332 [05:39<00:00,  1.02s/it]precision: 0.33
12 recall: 0.33
13 F1: 0.33
```

▼ Discussion

▼ Goal 3: Compare the pros and cons of rule-based and neural approaches.

Compare the pros and cons of both approaches with relevant examples from your experiments above. Concerning the accuracy, which approach would you choose to be used in a product? Explain.

From what I observed doing both parts, the pros and the cons of each approach is as follows:

The precision, F1 and recall measurements are much better for the first approach, for my assumption because we manually define all the rules which have semantic meaning and correctly and deterministically translating it to the SQL which gives us the desirable result.

Thus, working with a huge examples set and manually fitting all the expressions to deterministic translations we perform very good for those we've seen int he past. And if for exmaple we have a lot of popular simple requests (quries), the system in the first approach will perform very good. Those are the pros. It is also understanable.

The cons are that for an unseen examples (words of the request) we'll get bad performance, thus a system using that approach need to be handly supported and updated manually for the changing environment and external factors such as popularity of linguistic expressions and language place dependennt aspects and e.t.c

Another con is it takes a lot of proggraming effort and time to write manually all the derivations of the possible rules, another aspect of manual support needed for that approach.

In conclusion that approach can be good as we've seen, but need more handy work and time,with further support, which is also manual.

In contrast, the neaural approach, can be automatic and involves some programming only in writing the model, depsipte the fact it takes a lot of data and time to train, after that it does everything by itself, it also handles better unseen examples, and with a stronger model maybe could perform better then we saw (we didn't use a transformer with newest improvements or even a model like GPT-3). The adjustment to new changes, assume the language narrative and slang will change in a 20 years, the model can be adjusted more automatically, thought it also take time to retrain some of the model parameters and e.t.c. It less understandale then the previous approach, although with attention we can visualise it and see some semantic aspects of the way the model works (the attention and by it take a little glimpse into the model).

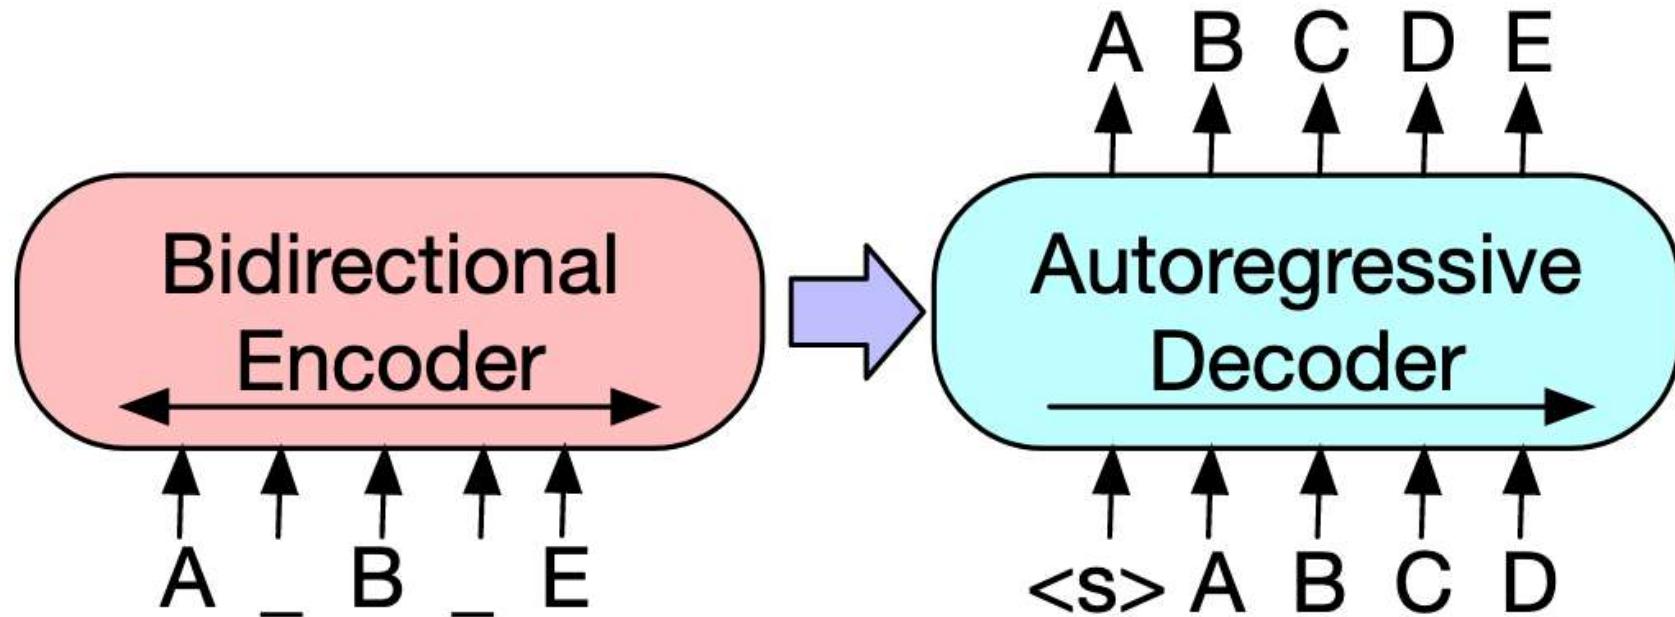
Probably in that setting and performance comparasion, I would use the first approach, but of course if the second approach, which takes more computation power, can achive similiar very good semnatic derivations, with huge datasets and e.t.c (like GPT-3 and e.t.c), then the automation and the benifits of that approach are at least invitint for further reasearch.

▼ (Optional) Goal 4: Use state-of-the-art pretrained transformers

The most recent breakthrough in Natural Language Processing stems from the usage of pretrained transformer models. For example, you might have heard of pretrained transformers such as [GPT-3](#) and [BERT](#)). These models are usually trained on vast amounts of text data using variants of language modeling objectives, and researchers have found that finetuning them on downstream tasks usually results in better performance as compared to training a model from scratch.

In the previous part, you implemented an LSTM-based sequence-to-sequence approach. To "upgrade" the model to be a state-of-the-art pretrained transformer only requires minor modifications.

The pretrained model that we will use is [BART](#), which uses a bidirectional transformer encoder and a unidirectional transformer decoder, as illustrated in the below diagram (image courtesy <https://arxiv.org/pdf/1910.13461.pdf>):



We can see that this model is strikingly similar to the LSTM-based encoder-decoder model we've been using. The only difference is that they use transformers instead of LSTMs. Therefore, we only need to change the modeling parts of the code, as we will see later.

First, we download and load the pretrained BART model from the [transformers](#) package by Huggingface. Note that we also need to use the "tokenizer" of BART, which is actually a combination of a tokenizer and a mapping from strings to word ids.

```
1 pretrained_bart = BartForConditionalGeneration.from_pretrained('facebook/bart-base')
2 bart_tokenizer = BartTokenizer.from_pretrained('facebook/bart-base')
```

Downloading: 100%	1.55k/1.55k [00:00<00:00, 3.58kB/s]
Downloading: 100%	558M/558M [00:17<00:00, 32.7MB/s]
Downloading: 100%	899k/899k [00:01<00:00, 791kB/s]
Downloading: 100%	456k/456k [00:00<00:00, 1.25MB/s]

Below we demonstrate how to use BART's tokenizer to convert a sentence to a list of word ids, and vice versa.

```

1 # BART uses a predefined "tokenizer", which directly maps a sentence
2 # to a list of ids
3 def bart_tokenize(string):
4     return bart_tokenizer(string)['input_ids'][:1024] # BART model can process at most 1024 tokens
5
6 def bart_detokenize(token_ids):
7     return bart_tokenizer.decode(token_ids, skip_special_tokens=True)
8
9 ## Demonstrating the tokenizer
10 question = 'Are there any first-class flights from St. Louis at 11pm for less than $3.50?'
11
12 tokenized_question = bart_tokenize(question)
13 print('tokenized:', tokenized_question)
14
15 detokenized_question = bart_detokenize(tokenized_question)
16 print('detokenized:', detokenized_question)

tokenized: [0, 13755, 89, 143, 78, 12, 4684, 4871, 31, 312, 4, 3217, 23, 365, 1685, 13, 540, 87, 68, 246, 4, 1096, 116
detokenized: Are there any first-class flights from St. Louis at 11pm for less than $3.50?

```

We need to reprocess the data using our new tokenizer. Note that here we set `batch_first` to `True`, since that's the expected input shape of the transformers package.

```
1 SRC_BART = tt.data.Field(include_lengths=True,      # include lengths
2                             batch_first=True,        # batches will be batch_size x max_len
3                             tokenize=bart_tokenize, # use bart tokenizer
4                             use_vocab=False,        # bart tokenizer already converts to int ids
5                             pad_token=bart_tokenizer.pad_token_id
6                         )
7 TGT_BART = tt.data.Field(include_lengths=False,
8                             batch_first=True,        # batches will be batch_size x max_len
9                             tokenize=bart_tokenize, # use bart tokenizer
10                            use_vocab=False,        # bart tokenizer already converts to int ids
11                            pad_token=bart_tokenizer.pad_token_id
12                         )
13 fields_bart = [('src', SRC_BART), ('tgt', TGT_BART)]
14
15 # Make splits for data
16 train_data_bart, val_data_bart, test_data_bart = tt.datasets.TranslationDataset.splits(
17     ('_flightid.nl', '_flightid.sql'), fields_bart, path='./data/',
18     train='train', validation='dev', test='test')
19
20 BATCH_SIZE = 1 # batch size for training/validation
21 TEST_BATCH_SIZE = 1 # batch size for test, we use 1 to make beam search implementation easier
22
23 train_iter_bart, val_iter_bart = tt.data.BucketIterator.splits((train_data_bart, val_data_bart),
24                                         batch_size=BATCH_SIZE,
25                                         device=device,
26                                         repeat=False,
27                                         sort_key=lambda x: len(x.src),
28                                         sort_within_batch=True)
29 test_iter_bart = tt.data.BucketIterator(test_data_bart,
30                                         batch_size=1,
31                                         device=device,
32                                         repeat=False,
33                                         sort=False,
34                                         train=False)
```

Token indices sequence length is longer than the specified maximum sequence length for this model (1135 > 1024). Runni

Let's take a look at the batch. Note that the shape of the batch is `batch_size x max_len`, instead of `max_len x batch_size` as in the previous part.

```

1 batch = next(iter(train_iter_bart))
2 train_batch_text, train_batch_text_lengths = batch.src
3 print (f"Size of text batch: {train_batch_text.shape}")
4 print (f"First sentence in batch: {train_batch_text[0]}")
5 print (f"Length of the third sentence in batch: {train_batch_text_lengths[0]}")
6 print (f"Converted back to string: {bart_detokenize(train_batch_text[0])}")
7
8 train_batch_sql = batch.tgt
9 print (f"Size of sql batch: {train_batch_sql.shape}")
10 print (f"First sql in batch: {train_batch_sql[0]}")
11 print (f"Converted back to string: {bart_detokenize(train_batch_sql[0])}")

Size of text batch: torch.Size([1, 16])
First sentence in batch: tensor([    0, 12196,     16,    110, 13342,    662,   2524,     31,  9473,    811,
    20048,      7, 16224,    462, 17034,      2], device='cuda:0')
Length of the third sentence in batch: 16
Converted back to string: what is your earliest morning flight from indianapolis to charlotte
Size of sql batch: torch.Size([1, 370])
First sql in batch: tensor([    0, 49179,    211, 11595,    2444,   7164,   2524,   1215,   134,      4,
    15801,   1215,    808, 11974,    2524,    2524,   1215,   134,   2156,   3062,
    1215, 11131,   3062,   1215, 11131,   1215,   134,   2156,   343,   343,
    1215,   134,   2156,   3062,   1215, 11131,   3062,   1215, 11131,   1215,
    176,   2156,   343,   343,   1215,   176, 29919,   2524,   1215,   134,
      4, 17272, 2013,   2407,   1215,    958,   5457,     36, 44664, 18335,
     36,   2524,   1215,   134,      4, 17272, 2013,   2407,   1215,    958,
    4839, 11974,   2524,   2524,   1215,   134,   2156,   3062,   1215, 11131,
    3062,   1215, 11131,   1215,   134,   2156,   343,   343,   1215,   134,
    2156,   3062,   1215, 11131,   3062,   1215, 11131,   1215,   176,   2156,
     343,   343,   1215,   176, 29919,   2524,   1215,   134,      4, 7761,
    1215,   2456, 3427,   5457,   3062,   1215, 11131,   1215,   134,      4,
    2456, 3427,   1215, 20414,   4248,   3062,   1215, 11131,   1215,   134,
      4, 14853,   1215, 20414,   5457,   343,   1215,   134,      4, 14853,
    1215, 20414,   4248,   343,   1215,   134,      4, 14853,   1215, 13650,
    5457,   128, 13796, 10296,   591,   3384,   1729,   108,   4248,     36,
    2524,   1215,   134,      4,   560,   1215,   2456,   3427,   5457,   3062,
    1215, 11131,   1215,   176,      4, 2456,   3427,   1215, 20414,   4248,
    3062,   1215, 11131,   1215,   176,      4, 14853,   1215, 20414,   5457,
```

```

343, 1215, 176, 4, 14853, 1215, 20414, 4248, 343, 1215,
176, 4, 14853, 1215, 13650, 5457, 128, 29146, 574, 3293,
6433, 108, 4248, 36, 2524, 1215, 134, 4, 17272, 2013,
2407, 1215, 958, 24844, 9112, 2796, 321, 4248, 23777, 4248,
112, 5457, 112, 4839, 4839, 4839, 4248, 36, 2524, 1215,
134, 4, 7761, 1215, 2456, 3427, 5457, 3062, 1215, 11131,
1215, 134, 4, 2456, 3427, 1215, 20414, 4248, 3062, 1215,
11131, 1215, 134, 4, 14853, 1215, 20414, 5457, 343, 1215,
134, 4, 14853, 1215, 20414, 4248, 343, 1215, 134, 4,
14853, 1215, 13650, 5457, 128, 13796, 10296, 591, 3384, 1729,
108, 4248, 36, 2524, 1215, 134, 4, 560, 1215, 2456,
3427, 5457, 3062, 1215, 11131, 1215, 176, 4, 2456, 3427,
1215, 20414, 4248, 3062, 1215, 11131, 1215, 176, 4, 14853,
1215, 20414, 5457, 343, 1215, 176, 4, 14853, 1215, 20414,
4248, 343, 1215, 176, 4, 14853, 1215, 13650, 5457, 128,
29146, 574, 3293, 6433, 108, 4248, 36, 2524, 1215, 134,
4, 17272, 2013, 2407, 1215, 958, 24844, 9112, 2796, 321,
4248, 23777, 4248, 112, 5457, 112, 4839, 4839, 4839, 4839, 2],
device='cuda:0')

```

Converted back to string: SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1,

Now we are ready to implement the BART-based approach for the text-to-SQL conversion problem. In the below `BART` class, we have provided the constructor `__init__`, the `forward` function, and the `predict` function. Your job is to implement the main optimization `train_all`, and `evaluate_ppl` for evaluating validation perplexity for model selection.

Hint: you can use almost the same `train_all` and `evaluate_ppl` function you implemented before, but here a major difference is that due to setting `batch_first=True`, the batched source/target tensors are of size `batch_size x max_len`, as opposed to `max_len x batch_size` in the LSTM-based approach, and you need to make changes in `train_all` and `evaluate_ppl` accordingly.

```

1 #TODO - finish implementing the `BART` class.
2 class BART(nn.Module):
3     def __init__(self, tokenizer, pretrained_bart):
4         """
5             Initializer. Creates network modules and loss function.
6             Arguments:
7                 tokenizer: BART tokenizer

```

```
    TOKENIZER = BART TOKENIZER
8     pretrained_bart: pretrained BART
9 """
10    super(BART, self).__init__()
11
12    self.V_tgt = len(tokenizer)
13
14    # Get special word ids
15    self.padding_id_tgt = tokenizer.pad_token_id
16
17    # Create essential modules
18    self.bart = pretrained_bart
19
20    # Create loss function
21    self.loss_function = nn.CrossEntropyLoss(reduction="sum",
22                                              ignore_index=self.padding_id_tgt)
23
24    def forward(self, src, src_lengths, tgt_in):
25        """
26            Performs forward computation, returns logits.
27            Arguments:
28                src: src batch of size (batch_size, max_src_len)
29                src_lengths: src lengths of size (batch_size)
30                tgt_in: a tensor of size (tgt_len, bsz)
31        """
32
33        # BART assumes inputs to be batch-first
34        # This single function is forwarding both encoder and decoder (w/ cross attn),
35        # using `input_ids` as encoder inputs, and `decoder_input_ids`
36        # as decoder inputs.
37        logots = self.bart(input_ids=src,
38                            decoder_input_ids=tgt_in,
39                            use_cache=False
40                            ).logots
41
42    return logots
43
44    def evaluate_ppl(self, iterator):
45        """Returns the model's perplexity on a given dataset `iterator`."""
46        # Switch to eval mode
47        self.eval()
48        +--+> 1000 - 0
```

```
40 total_loss = 0
41 total_words = 0
42 for batch in iterator:
43     # Input and target
44     src, src_lengths = batch.src # text: max_src_length, bsz
45     tgt = batch.tgt # max_tgt_length, bsz
46     tgt_in = tgt[0][:-1].unsqueeze(0) # Remove <eos> for decode input (y_0=<bos>, y_1, y_2)
47     tgt_out = tgt[0][1:].unsqueeze(0) # Remove <bos> as target (y_1, y_2, y_3=<eos>)
48     bsz = tgt.size(1) # remove <bos> as target (y_1, y_2, y_3=<eos>)
49     # Forward to get logits
50     logits = self.forward(src, src_lengths, tgt_in)
51     # Compute cross entropy loss
52     loss = self.loss_function(logits.squeeze(), tgt_out.squeeze())
53     total_loss += loss.item()
54     total_words += tgt_out.ne(self.padding_id_tgt).float().sum().item()
55
56 return math.exp(total_loss / total_words)
57
58
59
60
61
62
63 def train_all(self, train_iter, val_iter, epochs=10, learning_rate=0.001):
64     """Train the model."""
65     # Switch the module to training mode
66     self.train()
67     # Use Adam to optimize the parameters
68     optim = torch.optim.Adam(self.parameters(), lr=learning_rate)
69     best_validation_ppl = float('inf')
70     best_model = None
71     # Run the optimization for multiple epochs
72     for epoch in range(epochs):
73         total_words = 0
74         total_loss = 0.0
75         for batch in tqdm(train_iter):
76             # Zero the parameter gradients
77             self.zero_grad()
78             # Input and target
79             src, src_lengths = batch.src # text: max_src_length, bsz
80             tgt = batch.tgt # max_tgt_length, bsz
81             tgt_in = tgt[0][:-1].unsqueeze(0) # Remove <eos> for decode input (y_0=<bos>, y_1, y_2)
82             tgt_out = tgt[0][1:].unsqueeze(0) # Remove <bos> as target (y_1, y_2, y_3=<eos>)
83             bsz = tgt.size(1)
84             # Run forward pass and compute loss along the way.
85             # Add code here to compute loss and backpropagate
```

```
85     logits = self.forward(src, src_lengths, tgt_in)
86     loss = self.loss_function(logits.squeeze(), tgt_out.squeeze())
87     # Training stats
88     num_tgt_words = tgt_out.ne(self.padding_id_tgt).float().sum().item()
89     total_words += num_tgt_words
90     total_loss += loss.item()
91     # Perform backpropagation
92     loss.div(bsz).backward()
93     optim.step()
94
95     # Evaluate and track improvements on the validation dataset
96     validation_ppl = self.evaluate_ppl(val_iter)
97     self.train()
98     if validation_ppl < best_validation_ppl:
99         best_validation_ppl = validation_ppl
100        self.best_model = copy.deepcopy(self.state_dict())
101    epoch_loss = total_loss / total_words
102    print(f'Epoch: {epoch} Training Perplexity: {math.exp(epoch_loss):.4f} '
103          f'Validation Perplexity: {validation_ppl:.4f}')
104
105 def predict(self, tokens, K=1, max_T=400):
106     """
107     Generates the target sequence given the source sequence using beam search decoding.
108     Note that for simplicity, we only use batch size 1.
109     Arguments:
110         tokens: a list of strings, the source sentence.
111         max_T: at most proceed this many steps of decoding
112     Returns:
113         a string of the generated target sentence.
114     """
115     string = ' '.join(tokens) # first convert to a string
116     # Tokenize and map to a list of word ids
117     inputs = torch.LongTensor(bart_tokenize(string)).to(device).view(1, -1)
118     # The `transformers` package provides built-in beam search support
119     prediction = self.bart.generate(inputs,
120                                     num_beams=K,
121                                     max_length=max_T,
122                                     early_stopping=True,
123                                     no_repeat_ngram_size=0,
```

```

124         decoder_start_token_id=0,
125         use_cache=True)[0]
126
127     return bart_detokenize(prediction)

```

The code below will kick off training, and evaluate the validation perplexity. You should expect to see a value very close to 1.

```

1 EPOCHS = 5 # epochs, we recommend starting with a smaller number like 1
2 LEARNING_RATE = 1e-5 # learning rate
3
4 # Instantiate and train classifier
5 bart_model = BART(bart_tokenizer,
6                     pretrained_bart
7 ).to(device)
8
9 bart_model.train_all(train_iter_bart, val_iter_bart, epochs=EPOCHS, learning_rate=LEARNING_RATE)
10 bart_model.load_state_dict(bart_model.best_model)
11
12 # Evaluate model performance, the expected value should be < 1.2
13 print (f'Validation perplexity: {bart_model.evaluate_ppl(val_iter_bart):.3f}')

```

```

100%|██████████| 3651/3651 [07:31<00:00,  8.08it/s]
  0%|          | 1/3651 [00:00<06:40,  9.12it/s]Epoch: 0 Training Perplexity: 1.3980 Validation Perplexity: 1.0846
100%|██████████| 3651/3651 [07:31<00:00,  8.09it/s]
  0%|          | 1/3651 [00:00<06:41,  9.10it/s]Epoch: 1 Training Perplexity: 1.0916 Validation Perplexity: 1.0478
100%|██████████| 3651/3651 [07:31<00:00,  8.09it/s]
  0%|          | 1/3651 [00:00<09:46,  6.22it/s]Epoch: 2 Training Perplexity: 1.0565 Validation Perplexity: 1.0364
100%|██████████| 3651/3651 [07:31<00:00,  8.08it/s]
  0%|          | 0/3651 [00:00<?, ?it/s]Epoch: 3 Training Perplexity: 1.0399 Validation Perplexity: 1.0286
100%|██████████| 3651/3651 [07:30<00:00,  8.10it/s]
Epoch: 4 Training Perplexity: 1.0302 Validation Perplexity: 1.0243
Validation perplexity: 1.024

```

As before, make sure that your model is making reasonable predictions on a few examples before evaluating on the entire test set.

```

1 def bart_trial(sentence, gold_sql):
2     print("Sentence: ", sentence, "\n")
3     tokens = tokenize(sentence)

```

```
4  
5 predicted_sql = bart_model.predict(tokens, K=1, max_T=300)  
6 print("Predicted SQL:\n\n", predicted_sql, "\n")  
7  
8 if verify(predicted_sql, gold_sql, silent=False):  
9     print ('Correct!')  
10 else:  
11     print ('Incorrect!')
```

```
1 bart_trial(example_1, gold_sql_1)
```

Sentence: flights from phoenix to milwaukee

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1, city city_1, airport_serv
```

Predicted DB result:

```
[(108086,), (108087,), (301763,), (301764,), (301765,), (301766,), (302323,), (304881,), (310619,), (310620,)]
```

Gold DB result:

```
[(108086,), (108087,), (301763,), (301764,), (301765,), (301766,), (302323,), (304881,), (310619,), (310620,)]
```

Correct!

```
1 bart_trial(example_2, gold_sql_2)
```

Sentence: i would like a united flight

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 WHERE flight_1.airline_code = 'UA' AND flight_1.flight_number
```

Predicted DB result:

```
[]
```

Gold DB result:

```
[(100094,), (100099,), (100145,), (100158,), (100164,), (100167,), (100169,), (100203,), (100204,), (100296,)]
```

Incorrect!



```
1 bart_trial(example_3, gold_sql_3)
```

Sentence: i would like a flight between boston and dallas

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1, city city_1, airport_serv
```

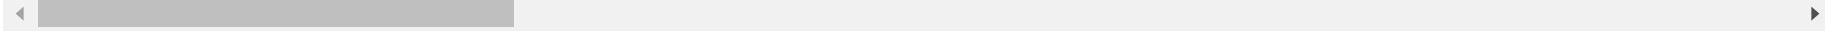
Predicted DB result:

```
[(103171,), (103172,), (103173,), (103174,), (103175,), (103176,), (103177,), (103178,), (103179,), (103180,)]
```

Gold DB result:

```
[(103171,), (103172,), (103173,), (103174,), (103175,), (103176,), (103177,), (103178,), (103179,), (103180,)]
```

Correct!



```
1 bart_trial(example_4, gold_sql_4)
```

Sentence: show me the united flights from denver to baltimore

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1, city city_1, airport_serv
```

Predicted DB result:

```
[(101231,), (101233,), (305983,)]
```

Gold DB result:

```
[(101231,), (101233,), (305983,)]
```

Correct!

```
1 bart_trial(example_5, gold_sql_5)
```

Sentence: show flights from cleveland to miami that arrive before 4pm

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1, city city_1, airport_serv
```

Predicted DB result:

```
[(107698,), (301117,)]
```

Gold DB result:

```
[(107698,), (301117,)]
```

Correct!

```
1 bart_trial(example_6, gold_sql_6b)
```

Sentence: okay how about a flight on sunday from tampa to charlotte

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1, city city_1, airport_serv
```

Predicted DB result:

```
[(101860,), (101861,), (101862,), (101863,), (101864,), (101865,), (305231,)]
```

Gold DB result:

```
[(101860,), (101861,), (101862,), (101863,), (101864,), (101865,), (305231,)]
```

Correct!

```
1 bart_trial(example_7, gold_sql_7b)
```

Sentence: list all flights going from boston to atlanta before 7 am on thursday

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1, city city_1, airport_serv
```

Predicted DB result:

```
[(100014,), (100015,), (100016,), (100017,), (100018,), (100019,), (304692,), (307330,), (100020,), (307329,)]
```

Gold DB result:

```
[(100014,)]
```

Incorrect!

```
1 bart_trial(example_8, gold_sql_8)
```

Sentence: list the flights from dallas to san francisco on american airlines

Predicted SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1, airport_service airport_service_1, city city_1, airport_serv
```

Predicted DB result:

```
[(108452,), (108454,), (108456,), (111083,), (111085,), (111086,), (111090,), (111091,), (111092,), (111094,)]
```

Gold DB result:

```
[(108452,), (108454,), (108456,), (111083,), (111085,), (111086,), (111090,), (111091,), (111092,), (111094,)]
```

Correct!

▼ Evaluation

The code below will evaluate on the entire test set. You should expect to see precision/recall/F1 greater than 40%.

```
1 def seq2seq_predictor_bart(tokens):
2     prediction = bart_model.predict(tokens, K=4, max_T=400)
3     return prediction
4
5
6 precision, recall, f1 = evaluate(seq2seq_predictor_bart, test_iter.dataset, num_examples=0)
7 print(f"precision: {precision:.2f}")
8 print(f"recall: {recall:.2f}")
9 print(f"F1: {f1:.2f}")
10
11 100%|██████████| 332/332 [22:07<00:00,  4.00s/it]precision: 0.50
12 recall: 0.50
13 F1: 0.50
```

▼ Debrief

Question: We're interested in any thoughts you have about this project segment so that we can improve it for later years, and to inform later segments for this year. Please list any issues that arose or comments you have to improve the project segment. Useful things to comment on include the following:

- Was the project segment clear or unclear? Which portions?
- Were the readings appropriate background for the project segment?
- Are there additions or changes you think would make the project segment better?

The project segment was clear and the reading and previous labs were appropriate and useful.

Instructions for submission of the project segment

This project segment should be submitted to Gradescope, which will be made available some time before the due date.

Project segment notebooks are manually graded, not autograded using otter as labs are. (Otter is used within project segment notebooks to synchronize distribution and solution code however.) **We will not run your notebook before grading it.** Instead, we ask that you submit the already freshly run notebook. The best method is to "restart kernel and run all cells", allowing time for all cells to be run to completion.

We also request that you **submit a PDF of the freshly run notebook**. The simplest method is to use "Export notebook to PDF", which will render the notebook to PDF via LaTeX. If that doesn't work, the method that seems to be most reliable is to export the notebook as HTML (if you are using Jupyter Notebook, you can do so using `File -> Print Preview`), open the HTML in a browser, and print it to a file. Then make sure to add the file to your git commit. Please name the file the same name as this notebook, but with a `.pdf` extension. (Conveniently, the methods just described will use that name by default.) You can then perform a git commit and push and submit the commit to Gradescope.

End of project segment 4

